

FCV Potential Customers Prediction

November 25, 2024

1 ExtraaLearn Project

1.1 Context

The EdTech industry has been surging in the past decade immensely, and according to a forecast, the Online Education market would be worth \$286.62bn by 2023 with a compound annual growth rate (CAGR) of 10.26% from 2018 to 2023. The modern era of online education has enforced a lot in its growth and expansion beyond any limit. Due to having many dominant features like ease of information sharing, personalized learning experience, transparency of assessment, etc, it is now preferable to traditional education.

In the present scenario due to the Covid-19, the online education sector has witnessed rapid growth and is attracting a lot of new customers. Due to this rapid growth, many new companies have emerged in this industry. With the availability and ease of use of digital marketing resources, companies can reach out to a wider audience with their offerings. The customers who show interest in these offerings are termed as leads. There are various sources of obtaining leads for Edtech companies, like

- The customer interacts with the marketing front on social media or other online platforms.
- The customer browses the website/app and downloads the brochure
- The customer connects through emails for more information.

The company then nurtures these leads and tries to convert them to paid customers. For this, the representative from the organization connects with the lead on call or through email to share further details.

1.2 Objective

ExtraaLearn is an initial stage startup that offers programs on cutting-edge technologies to students and professionals to help them upskill/reskill. With a large number of leads being generated on a regular basis, one of the issues faced by ExtraaLearn is to identify which of the leads are more likely to convert so that they can allocate resources accordingly. You, as a data scientist at ExtraaLearn, have been provided the leads data to:

- * Analyze and build an ML model to help identify which leads are more likely to convert to paid customers,
- * Find the factors driving the lead conversion process
- * Create a profile of the leads which are likely to convert

1.3 Data Description

The data contains the different attributes of leads and their interaction details with ExtraaLearn. The detailed data dictionary is given below.

Data Dictionary * ID: ID of the lead * age: Age of the lead * current_occupation: Current occupation of the lead. Values include 'Professional', 'Unemployed', and 'Student' * first_interaction: How did the lead first interacted with ExtraaLearn. Values include 'Website', 'Mobile App' * profile_completed: What percentage of profile has been filled by the lead on the website/mobile app. Values include Low - (0-50%), Medium - (50-75%), High (75-100%) * website_visits: How many times has a lead visited the website * time_spent_on_website: Total time spent on the website * page_views_per_visit: Average number of pages on the website viewed during the visits. * last_activity: Last interaction between the lead and ExtraaLearn. * Email Activity: Seeking for details about program through email, Representative shared information with lead like brochure of program, etc * Phone Activity: Had a Phone Conversation with representative, Had conversation over SMS with representative, etc * Website Activity: Interacted on live chat with representative, Updated profile on website, etc

- print_media_type1: Flag indicating whether the lead had seen the ad of ExtraaLearn in the Newspaper.
- print_media_type2: Flag indicating whether the lead had seen the ad of ExtraaLearn in the Magazine.
- digital_media: Flag indicating whether the lead had seen the ad of ExtraaLearn on the digital platforms.
- educational_channels: Flag indicating whether the lead had heard about ExtraaLearn in the education channels like online forums, discussion threads, educational websites, etc.
- referral: Flag indicating whether the lead had heard about ExtraaLearn through reference.
- status: Flag indicating whether the lead was converted to a paid customer or not.

1.4 Importing necessary libraries and data

```
[30]: import warnings

warnings.filterwarnings("ignore")
from statsmodels.tools.sm_exceptions import ConvergenceWarning

warnings.simplefilter("ignore", ConvergenceWarning)

# Libraries to help with reading and manipulating data

import pandas as pd
import numpy as np

# Library to split data
from sklearn.model_selection import train_test_split

# libraries to help with data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Removes the limit for the number of displayed columns
pd.set_option("display.max_columns", None)
# Sets the limit for the number of displayed rows
```

```

pd.set_option("display.max_rows", 200)
# setting the precision of floating numbers to 5 decimal points
pd.set_option("display.float_format", lambda x: "%.5f" % x)

# To build model for prediction
import statsmodels.stats.api as sms
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
from statsmodels.tools.tools import add_constant
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn import tree
from sklearn.ensemble import RandomForestClassifier

# To tune different models
from sklearn.model_selection import GridSearchCV

# To get diferent metric scores
import sklearn.metrics as metrics
from sklearn.metrics import (
    f1_score,
    accuracy_score,
    recall_score,
    precision_score,
    confusion_matrix,
    classification_report,
    roc_auc_score,
    precision_recall_curve,
    roc_curve,
    make_scorer,
)

```

1.5 Data Overview

- Let's see how the dataframe looks like

```

[31]: original_data = pd.read_csv(r"C:\Users\jprey\OneDrive\Escritorio\JP\KUL\5th_
↳Semester (Polimi)\MIT IDSS\Classification and Hypothesis_
↳Testing\Extraa_learn.csv")
data = original_data.copy()

```

```

[32]: data.head()

```

```

[32]:      ID  age  current_occupation  first_interaction  profile_completed \
0  EXT001   57        Unemployed        Website        High
1  EXT002   56    Professional    Mobile App        Medium

```

2	EXT003	52	Professional	Website	Medium
3	EXT004	53	Unemployed	Website	High
4	EXT005	23	Student	Website	High

	website_visits	time_spent_on_website	page_views_per_visit	\
0	7	1639	1.86100	
1	2	83	0.32000	
2	3	330	0.07400	
3	4	464	2.05700	
4	4	600	16.91400	

	last_activity	print_media_type1	print_media_type2	digital_media	\
0	Website Activity	Yes	No	Yes	
1	Website Activity	No	No	No	
2	Website Activity	No	No	Yes	
3	Website Activity	No	No	No	
4	Email Activity	No	No	No	

	educational_channels	referral	status
0	No	No	1
1	Yes	No	0
2	No	No	0
3	No	No	1
4	No	No	0

```
[33]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4612 entries, 0 to 4611
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   ID                    4612 non-null   object
1   age                   4612 non-null   int64
2   current_occupation    4612 non-null   object
3   first_interaction      4612 non-null   object
4   profile_completed      4612 non-null   object
5   website_visits         4612 non-null   int64
6   time_spent_on_website  4612 non-null   int64
7   page_views_per_visit   4612 non-null   float64
8   last_activity          4612 non-null   object
9   print_media_type1      4612 non-null   object
10  print_media_type2      4612 non-null   object
11  digital_media          4612 non-null   object
12  educational_channels    4612 non-null   object
13  referral               4612 non-null   object
14  status                 4612 non-null   int64
```

```
dtypes: float64(1), int64(4), object(10)
memory usage: 540.6+ KB
```

```
[34]: data.drop_duplicates(inplace=True)
```

1.5.1 Observations from Data Overview:

- 4612 datapoints (rows), 15 categories collected (columns)
- No Duplicates, no null values
- Only Age, Website Visits, Time Spent On Website, Page Views Per Visit and Status are numerical values. Although status is binary classification rather than a numerical datatype.

1.6 Exploratory Data Analysis (EDA)

1.7 Data Preprocessing

```
[35]: # The first column will not be helpful for our analysis.
data = data.drop("ID", axis=1)
data.head()
```

```
[35]:
```

	age	current_occupation	first_interaction	profile_completed	website_visits	\
0	57	Unemployed	Website	High	7	
1	56	Professional	Mobile App	Medium	2	
2	52	Professional	Website	Medium	3	
3	53	Unemployed	Website	High	4	
4	23	Student	Website	High	4	

	time_spent_on_website	page_views_per_visit	last_activity	\
0	1639	1.86100	Website Activity	
1	83	0.32000	Website Activity	
2	330	0.07400	Website Activity	
3	464	2.05700	Website Activity	
4	600	16.91400	Email Activity	

	print_media_type1	print_media_type2	digital_media	educational_channels	\
0	Yes	No	Yes	No	
1	No	No	No	Yes	
2	No	No	Yes	No	
3	No	No	No	No	
4	No	No	No	No	

	referral	status
0	No	1
1	No	0
2	No	0
3	No	1
4	No	0

```
[36]: # Making a list of all catrgorical variables
categorical_columns = list(data.select_dtypes("object").columns)
categorical_columns.append("status")

# Printing the different filled in values per categorical category
for column in categorical_columns:
    print(data[column].value_counts())
    print("-" * 50)
```

```
current_occupation
Professional    2616
Unemployed     1441
Student         555
Name: count, dtype: int64
-----
```

```
first_interaction
Website        2542
Mobile App     2070
Name: count, dtype: int64
-----
```

```
profile_completed
High           2264
Medium         2241
Low            107
Name: count, dtype: int64
-----
```

```
last_activity
Email Activity  2278
Phone Activity  1234
Website Activity 1100
Name: count, dtype: int64
-----
```

```
print_media_type1
No             4115
Yes            497
Name: count, dtype: int64
-----
```

```
print_media_type2
No             4379
Yes            233
Name: count, dtype: int64
-----
```

```
digital_media
No             4085
Yes            527
Name: count, dtype: int64
-----
```

```
educational_channels
No      3907
Yes      705
Name: count, dtype: int64
```

```
referral
No      4519
Yes      93
Name: count, dtype: int64
```

```
status
0      3235
1      1377
Name: count, dtype: int64
```

Observations:

- The data answers seems coherent suggesting a clean dataset.
- Each column has 2 or 3 answers at most. This categorical columns can be used to “hue” the different plots.

1.7.1 Univariate Analysis

Function to generate a Box Plot on top of a Histogram.

```
[ ]: # function to plot a boxplot and a histogram along the same scale for valuable
      ↳ visualization

def histogram_boxplot(data, feature, figsize=(12, 7), kde=False, bins=None):
    """
    Boxplot and histogram combined

    data: dataframe going to be used
    feature: dataframe column to be evaluated
    figsize: size of the figure (default (12,7))
    kde: whether to the show density curve (default False)
    bins: number of bins for histogram (default None)
    """

    # For the histogram
    if bins: # If bins parameter is passed
        sns.histplot(data=data, x=feature, kde=kde, bins=bins, palette="winter")
    else:
        sns.histplot(data=data, x=feature, kde=kde)
    # Vertical line for the mean
    plt.axvline(data[feature].mean(), color="black", linestyle="--")

    # Building a Box Plot with vertical dotted lines
```

```

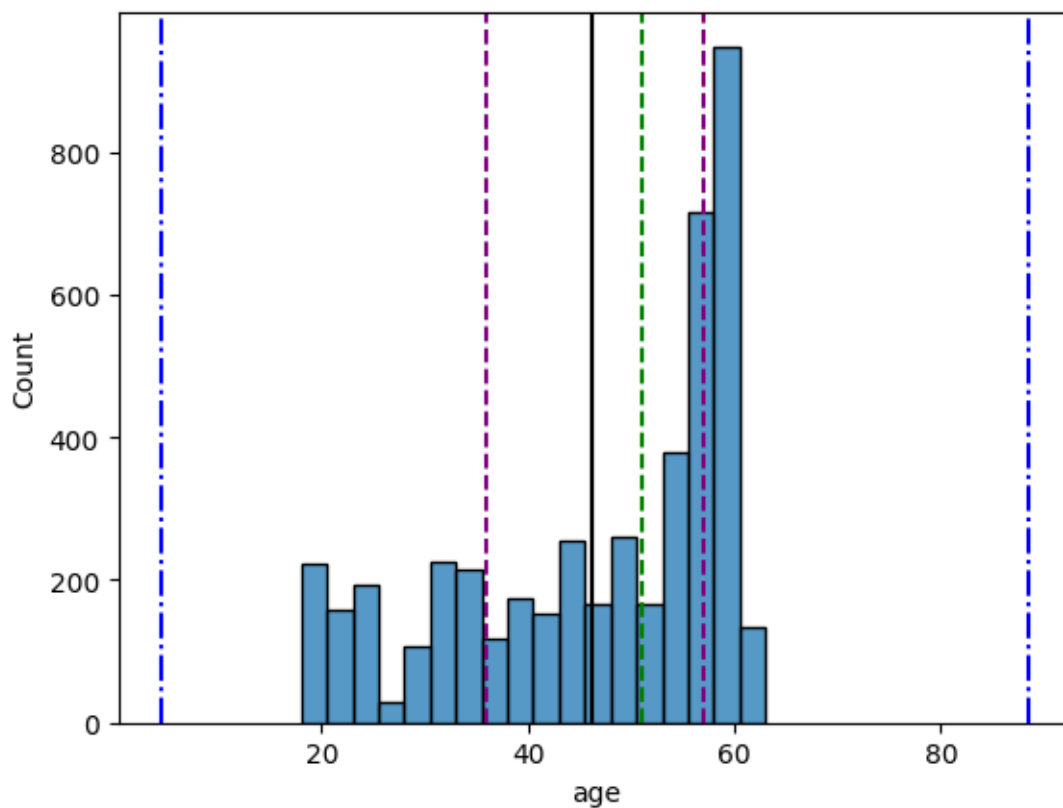
iqr = data[feature].quantile(0.75) - data[feature].quantile(0.25)

# Plot the "box plot"
plt.axvline(data[feature].median(), color="green", linestyle="--")
plt.axvline(data[feature].quantile(0.25), color="purple", linestyle="--")
plt.axvline(data[feature].quantile(0.75), color="purple", linestyle="--")
plt.axvline(data[feature].quantile(0.25) - 1.5*iqr, color="blue",
↪linestyle="-.")
plt.axvline(data[feature].quantile(0.75) + 1.5*iqr, color="blue",
↪linestyle="-.")
plt.show()

```

Let's explore the Age category

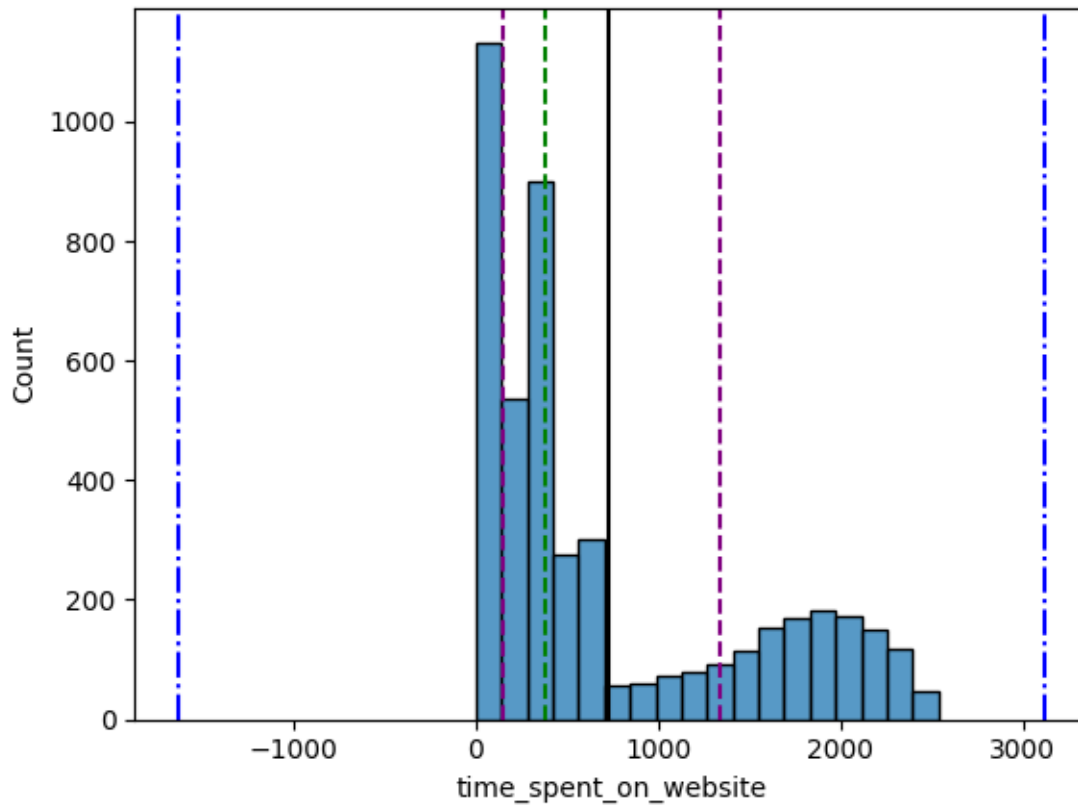
```
[38]: histogram_boxplot(data, "age")
```



Age values seem in order and no outlier treatment needed.

Let's explore the Time Spent on Website Category

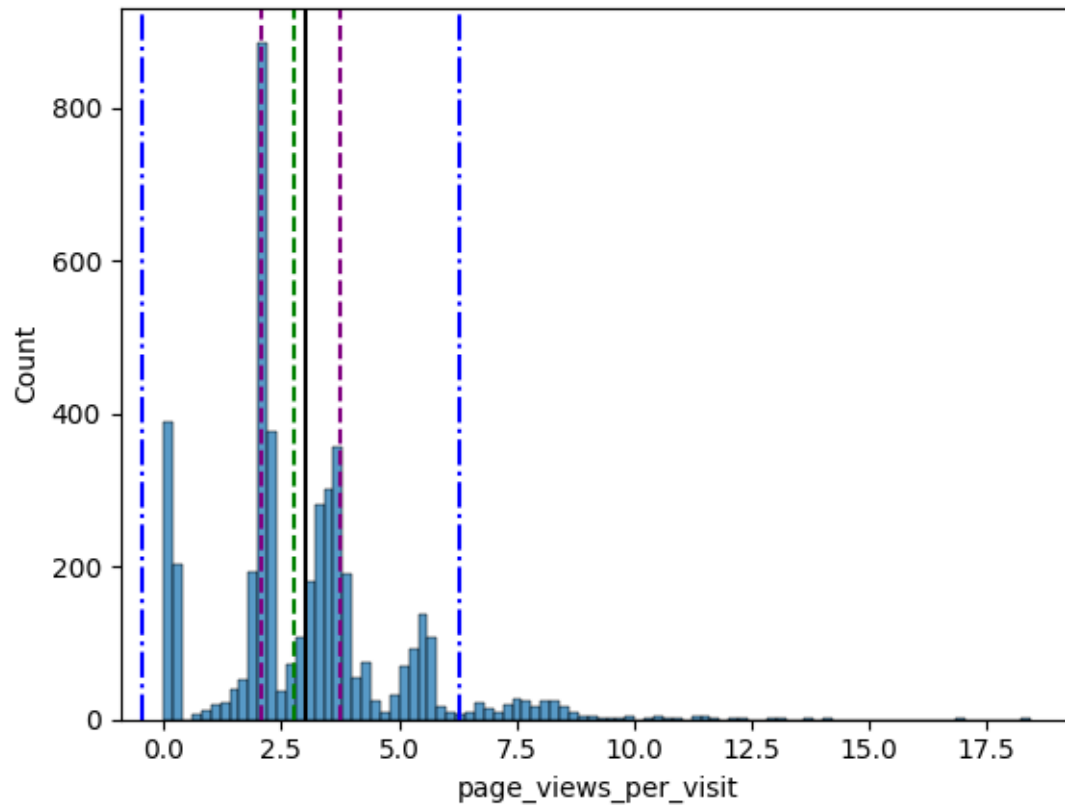
```
[39]: histogram_boxplot(data, "time_spent_on_website")
```

Time Spent on Website values seem in order and no outlier treatment needed. It is pertinent to mention that the distribution is rather bimodal.

Let's explore the Page Views per Visit Category

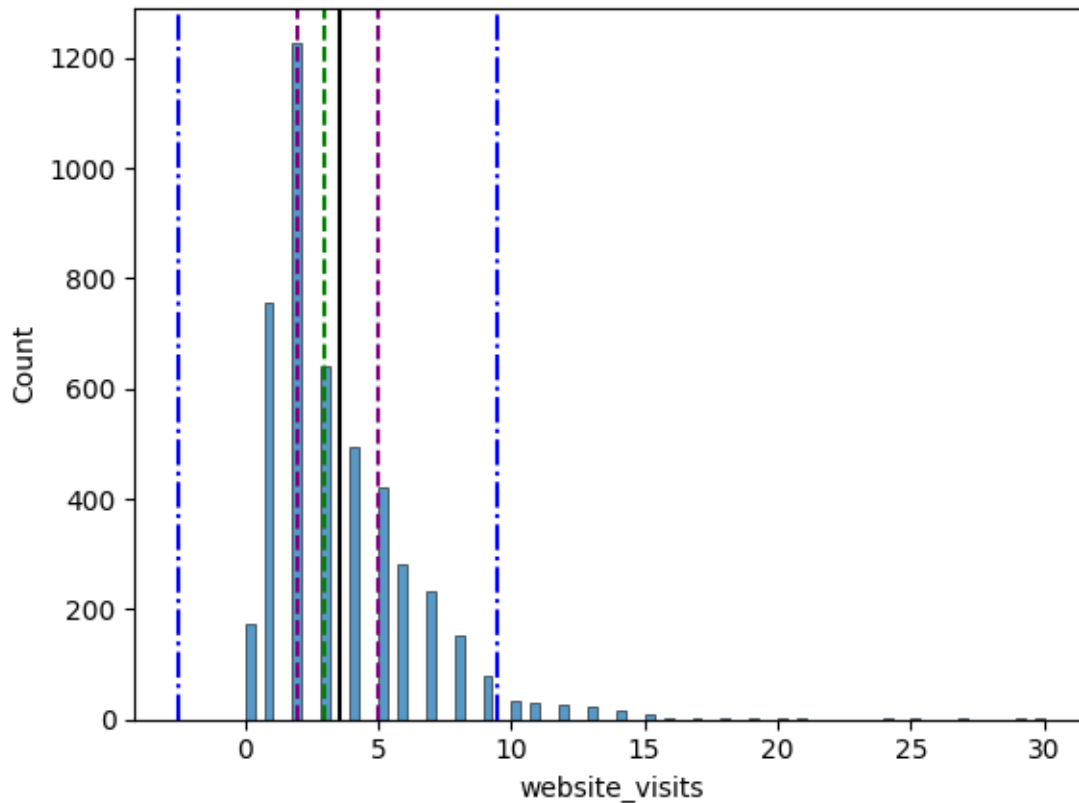
```
[40]: histogram_boxplot(data, "page_views_per_visit")
```



Page Views Per Visit values seem in order and no outlier treatment needed. The outliers come from some very frequent clients, which is understandable. The distribution is rather multimodal.

Let's consider the amount of Website Visits

```
[41]: histogram_boxplot(data, "website_visits")
```



The distribution is right skewed. There are some outliers, but yet again, it makes sense for some visitors to visit the website a lot. No outliers treatment needed.

Let's proceed with the univariate analysis for categorical variables.

Function for Categorical Columns. Visualization of the distribution of values

```
[82]: # Function to create labeled barplots for categorical columns
def labeled_barplot(data, feature, perc=False, n=None):
    """
    Barplot with percentage label at the top

    data: dataframe to be considered
    feature: existing column in the dataframe to be considered
    perc: whether to display percentages instead of count (default is False)
    n: displays the top n category levels (default is None, i.e., display all
    ↪ levels)
    """

    number_of_datapoints = len(data[feature]) # Total number of datapoints
    nr_of_categories = data[feature].nunique() # Counts the number of categories
```

```

if n is None:
    plt.figure(figsize=(nr_of_categories + 1, 5))
else:
    plt.figure(figsize=(n + 1, 5))

plt.xticks(rotation=90, fontsize=15)
ax = sns.countplot(
    data=data,
    x=feature,
    palette="hls",
    order=data[feature].value_counts().index[:n].sort_values(),
    # column from dataframe => count the occurrences per category =>...
    # ...=> get the first n elements (in n=None => index[:None]=index[:] =>
    ↪ it gets all)... =>
    # ...=> sort them by alphabetical order of the categories.
)
xtick_labels = [tick.get_text() for tick in ax.get_xticklabels()] # Store
    ↪ the name in a list
print("-"*5, feature, "-"*5)

for i, p in enumerate(ax.patches): # a patch is a bar <=> a category
    category_name = xtick_labels[i] # get the name
    if perc == True:
        # Formatted to one decimal place and percentage symbol included.
        # get_height() returns the height of the bar <=> instances of that
    ↪ category
        label = "{:.1f}%".format(100 * p.get_height() /
    ↪ number_of_datapoints)
    else:
        # get_height() returns the height of the bar <=> instances of that
    ↪ category
        label = p.get_height()
        print(f"{category_name}: {round(100*p.get_height()/
    ↪ number_of_datapoints, 1)}%")

    # Where to position the label
    # Horizontally: find the beginning of the bar, and add half of its
    ↪ length (centered)
    x = p.get_x() + p.get_width() / 2
    y = p.get_height() # Vertically: Height of the bar

    # Writing the label
    ax.annotate(
        label,
        # The text to display (either percentage or count)

```

```

    (x, y),          # Coordinates (x, y) where the text should be
    ↪positioned
    ha="center",     # Horizontal alignment of the text relative to
    ↪the (x, y) point
    va="center",     # Vertical alignment of the text relative to the
    ↪(x, y) point
    size=12,         # Font size of the text
    xytext=(0, 5),   # Offset of the text from (x, y) in (x, y)
    ↪coordinates
    textcoords="offset points" # Interprets xytext as an offset in points
    ↪(not data units)
    )

plt.show() # show the plot

```

Let's visualize each categorical column

```

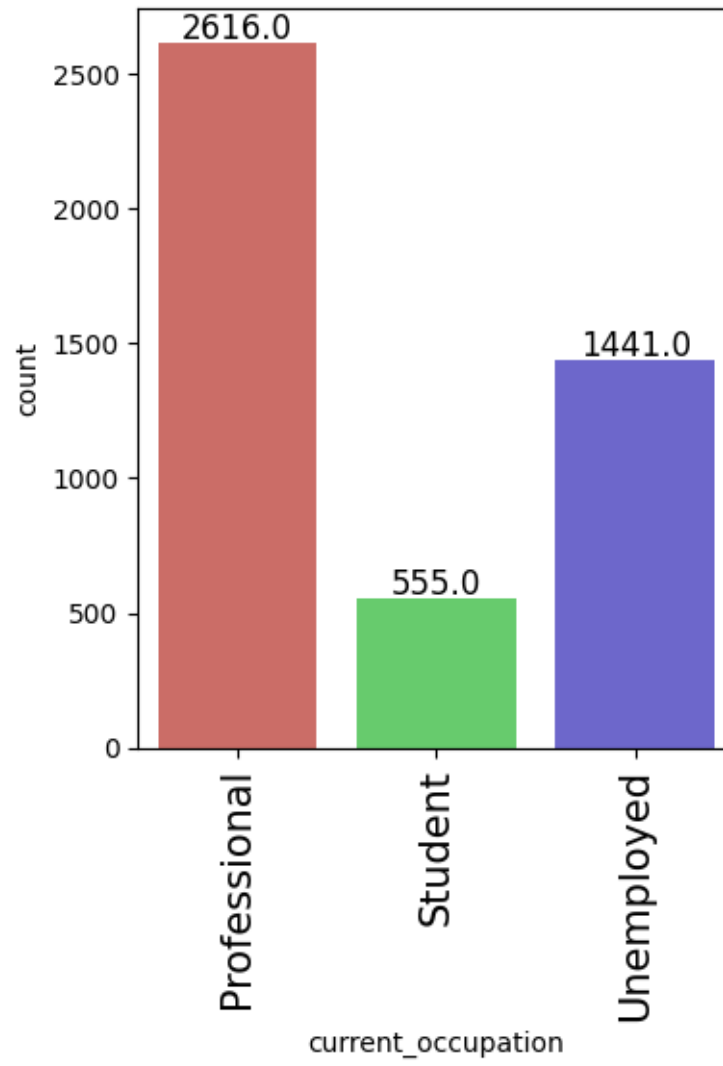
[83]: for category in categorical_columns:
        labeled_barplot(data, category)

```

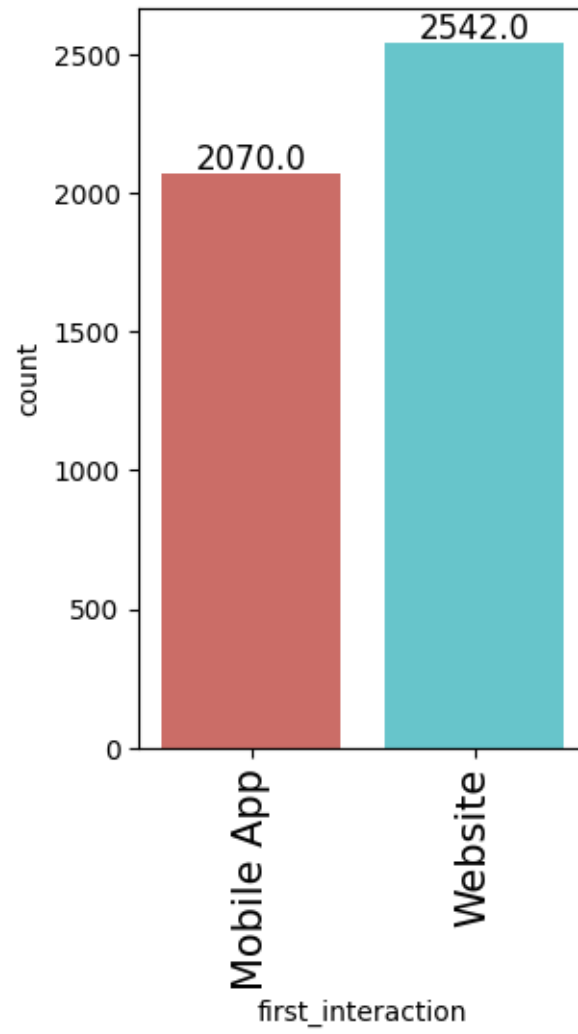
```

----- current_occupation -----
Professional: 56.7%
Student: 12.0%
Unemployed: 31.2%

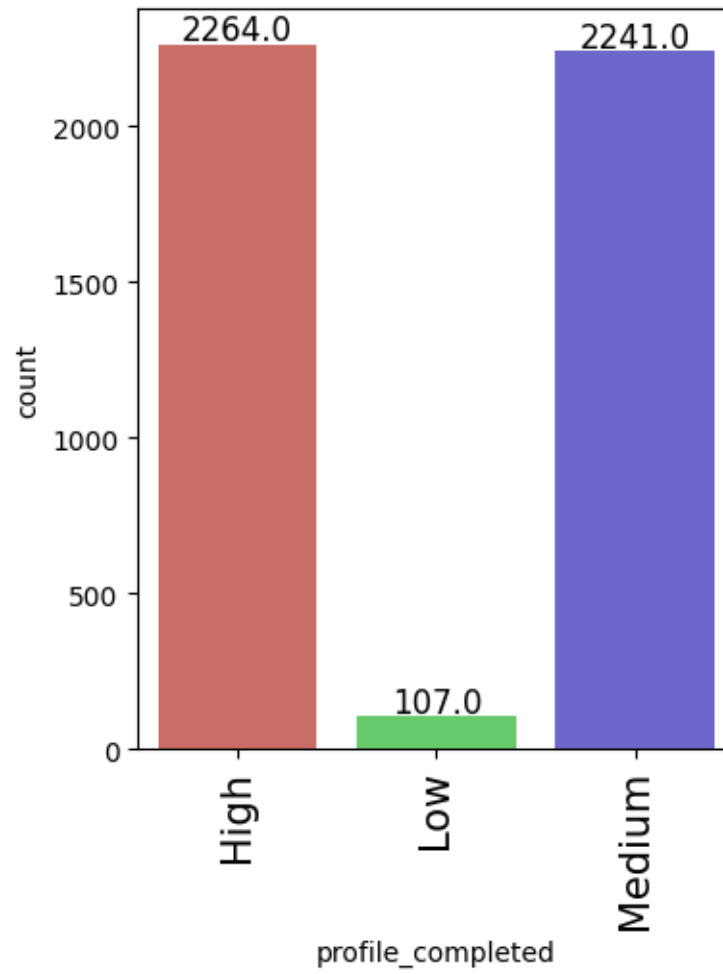
```



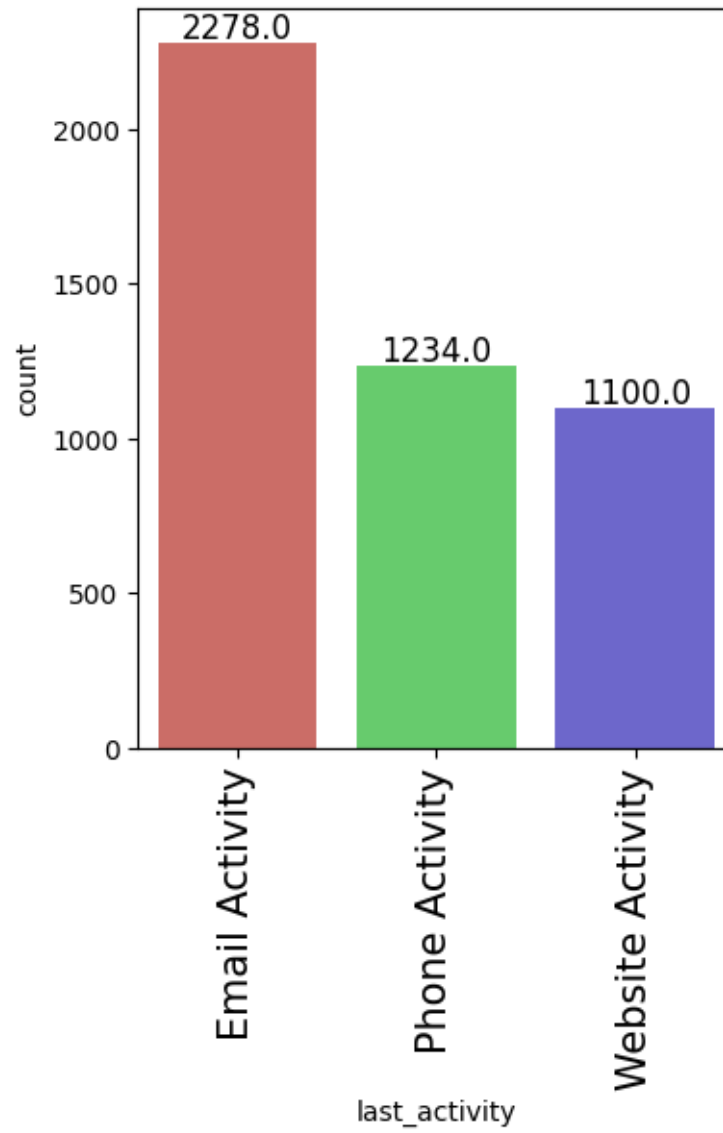
----- first_interaction -----
Mobile App: 44.9%
Website: 55.1%



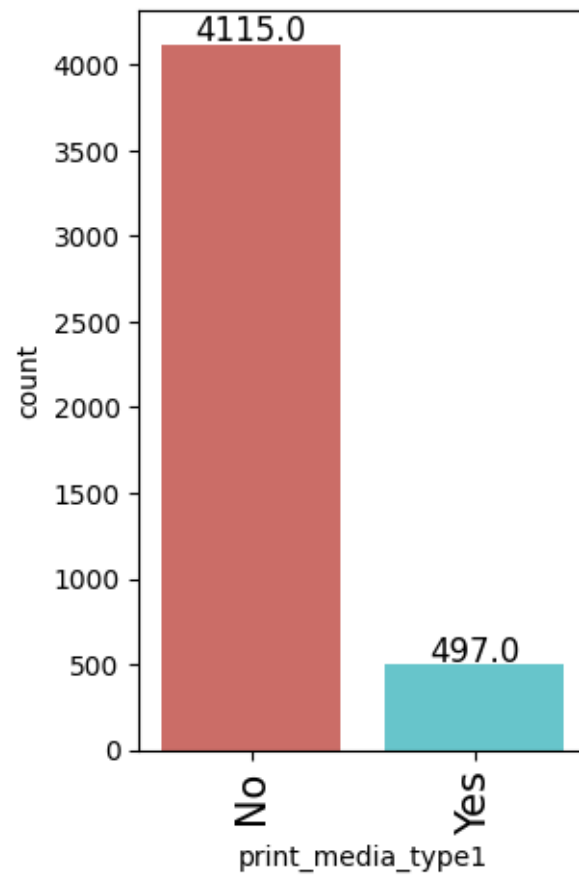
----- profile_completed -----
High: 49.1%
Low: 2.3%
Medium: 48.6%



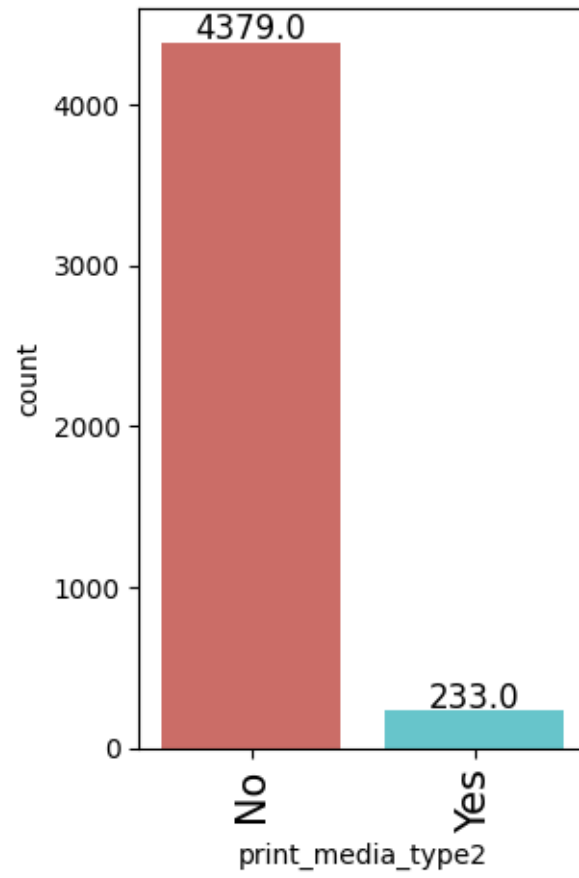
----- last_activity -----
Email Activity: 49.4%
Phone Activity: 26.8%
Website Activity: 23.9%



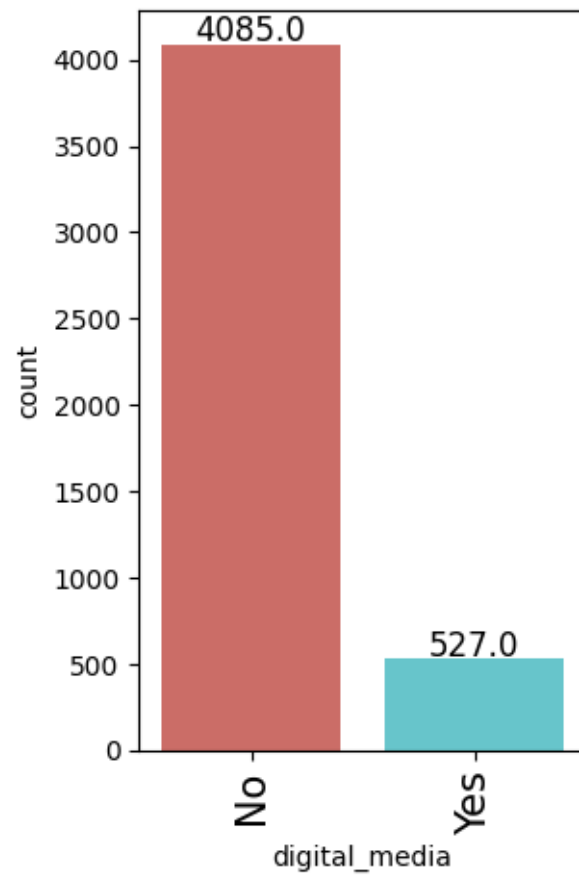
```
----- print_media_type1 -----  
No: 89.2%  
Yes: 10.8%
```



----- print_media_type2 -----
No: 94.9%
Yes: 5.1%



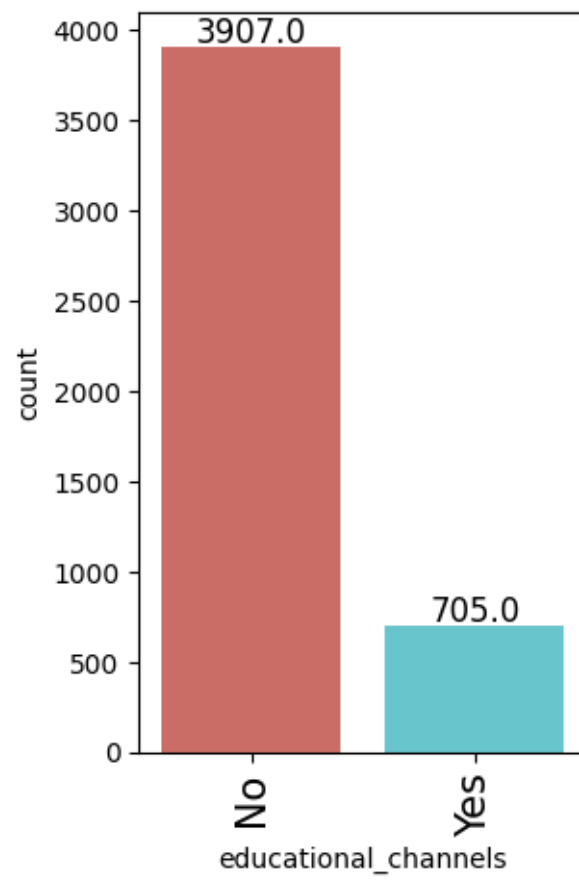
----- digital_media -----
No: 88.6%
Yes: 11.4%



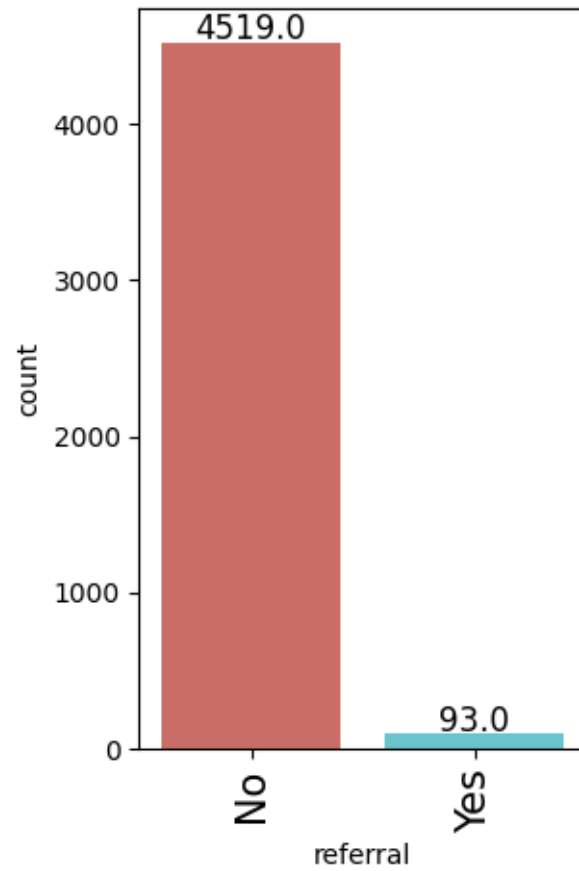
----- educational_channels -----

No: 84.7%

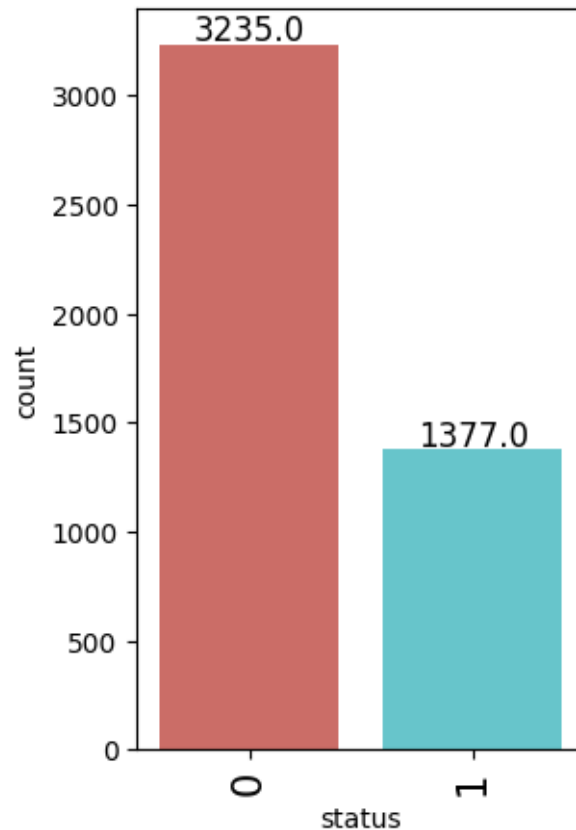
Yes: 15.3%



----- referral -----
No: 98.0%
Yes: 2.0%



----- status -----
0: 70.1%
1: 29.9%



Observations:

- 56% percent of interactions come from professionals. There is a strong number (31.2%) of interactions that come from unemployed people. It will be interesting to do a bivariate analysis considering the occupation and the status.
- Mobile App and Website are similarly used.
- Only 100 people have completed their profile less than 50%.
- Email interaction seems to have been the latest activity for most leads.
- Only 44.6% of leads heard/saw the program thanks to an ad/referral.
- Almost 30% of leads become paid customers.

1.7.2 Bivariate Analysis

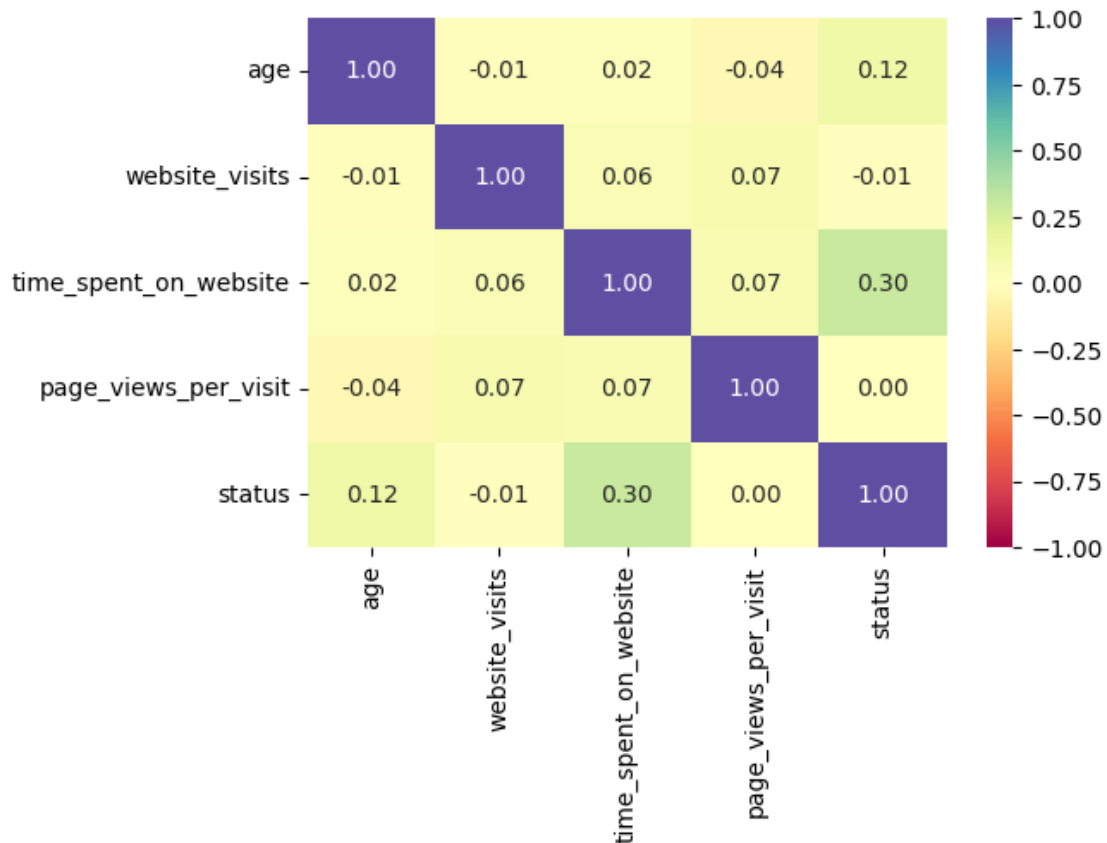
It is often helpful to begin the Bivariate Analysis with a Heatmap and a Pairplot

Heatmap

```
[44]: # Extract the numerical columns
cols_list = data.select_dtypes(include=np.number).columns.tolist()

# Create plot using Seaborn Heatmap
```

```
plt.figure(figsize=(6, 4))
sns.heatmap(data[cols_list].corr(), annot=True, vmin=-1, vmax=1, fmt=".2f",
            cmap="Spectral")
''' annot=True to visualize also the numbers
    vmin and vmax show the range of the correlation values
    fmt defines the numerical format (decimals)
    cmap is the colormap to be used.'''
plt.show()
```

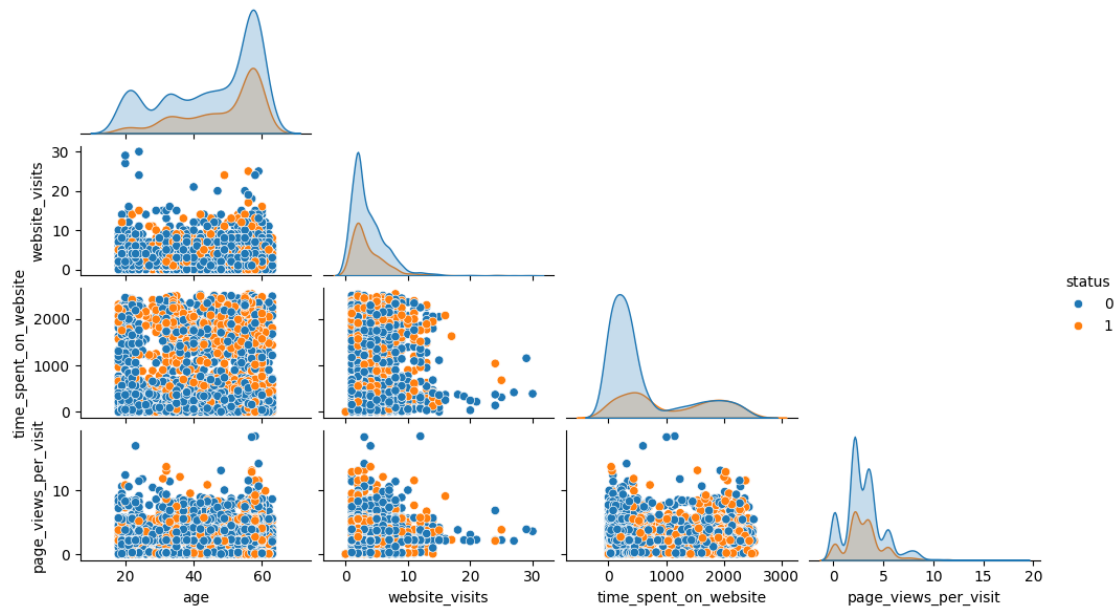


The heatmap shows that there is almost no relation between status and the rest of the categories. Time spent on Website shows a small correlation, which will be exploited later on.

Pairplot

```
[45]: plt.figure(figsize=(5, 4))
sns.pairplot(data, height=1.5, aspect=1.7, corner=True, diag_kind="kde",
            hue="status")
plt.show()
```


<Figure size 500x400 with 0 Axes>

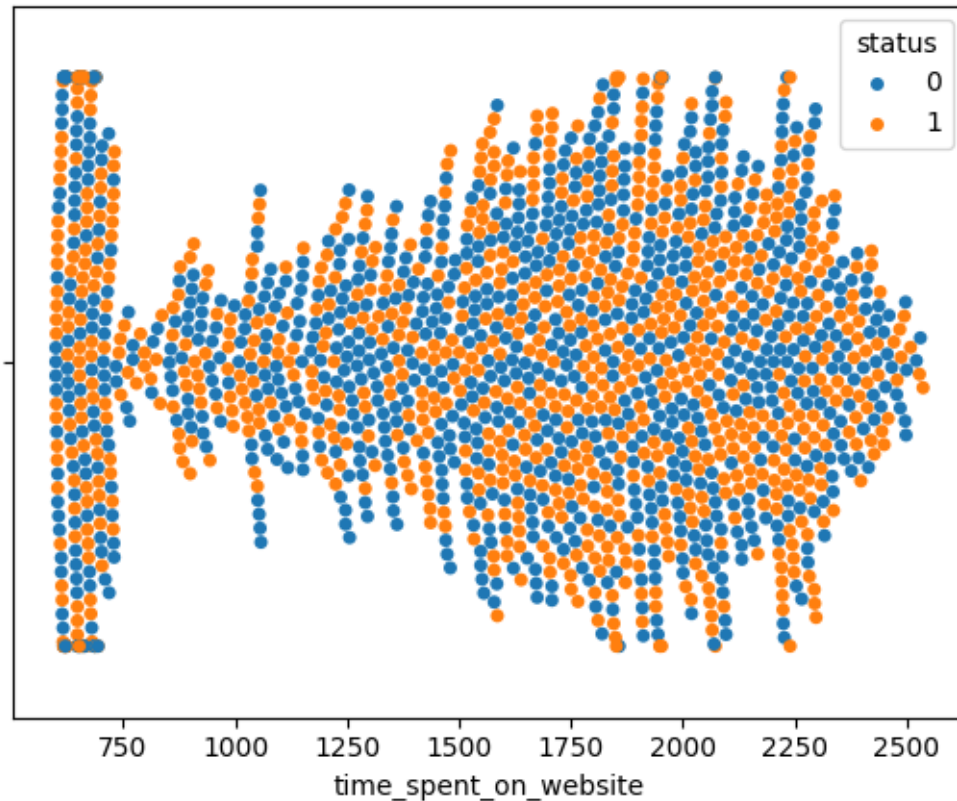


Although the pairplot did not illustrate as much as it does sometimes, it does illustrate that the time spent on website is clearly relevant for those with more than 600 minutes. Namely, it seems that 50% of those leads will become paid customers.

Let's explore this even more in depth with a swarmplot

```
[46]: high_time_on_website = data[data["time_spent_on_website"] >= 600]
      print(high_time_on_website["status"].value_counts())
      sns.swarmplot(data=high_time_on_website, x="time_spent_on_website",
                    hue="status")
      plt.show()
```

```
status
0      861
1      821
Name: count, dtype: int64
```



This is indeed interesting and will probably be one of the first questions in the decision tree.

Let's explore if other categories have an impact on whether the lead becomes a customer depending on the answers.

To do so, 3 helper functions will be defined.

[47]: *# Function that returns the order of the names of a dataframe.*

```
def get_order_categories(data):
    names = []
    for i in data.index:
        if (i[0] not in names):
            names.append(i[0])
    return names # Order of the names in this dataframe.
```

[48]: *# Creates a dataframe with proportion and count of a category per status*

```
def proportion_and_count(data, feature, hue="status"):
    data_proportion = data.groupby(feature)[hue].
    ↪ value_counts(normalize=True)*100
    data_count = data.groupby(feature)[hue].value_counts(normalize=False)
```

```

joint_dataframe = pd.concat([data_proportion, data_count], axis=1)
joint_dataframe = joint_dataframe.sort_values(by=[feature, hue])
return joint_dataframe, get_order_categories(joint_dataframe)

```

```

[49]: # Creates countplot and prints the relative/absolute quantities of a given
      ↪ category.
def barplot_hued(data, feature, hue="status", figsize=None):
    if (figsize):
        plt.figure(figsize=figsize)
    joint_dataframe, ordered_names = proportion_and_count(data, feature, hue)
    print(joint_dataframe)
    sns.countplot(data=data, x=feature, hue=hue, order=ordered_names)
    plt.title(f"{feature} shown per {hue}")
    plt.show()

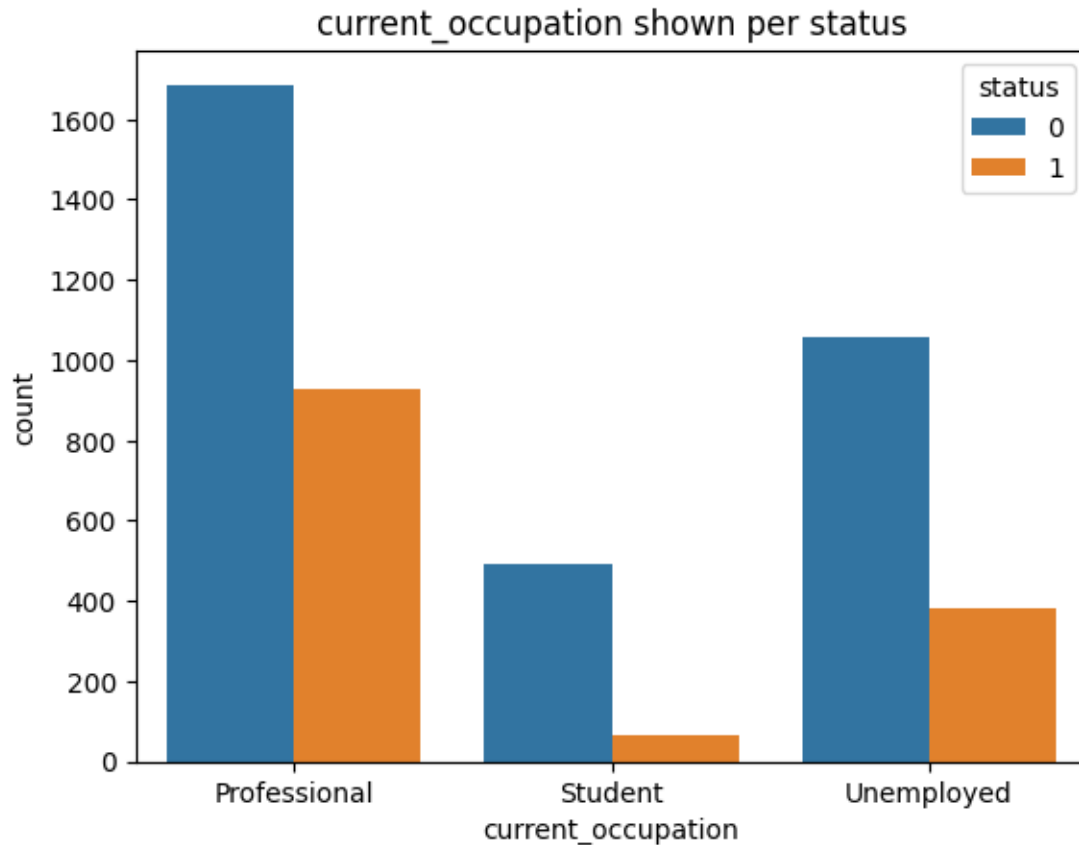
```

```

[50]: barplot_hued(data, "current_occupation", hue="status")

```

		proportion	count
current_occupation	status		
Professional	0	64.48777	1687
	1	35.51223	929
Student	0	88.28829	490
	1	11.71171	65
Unemployed	0	73.42124	1058
	1	26.57876	383

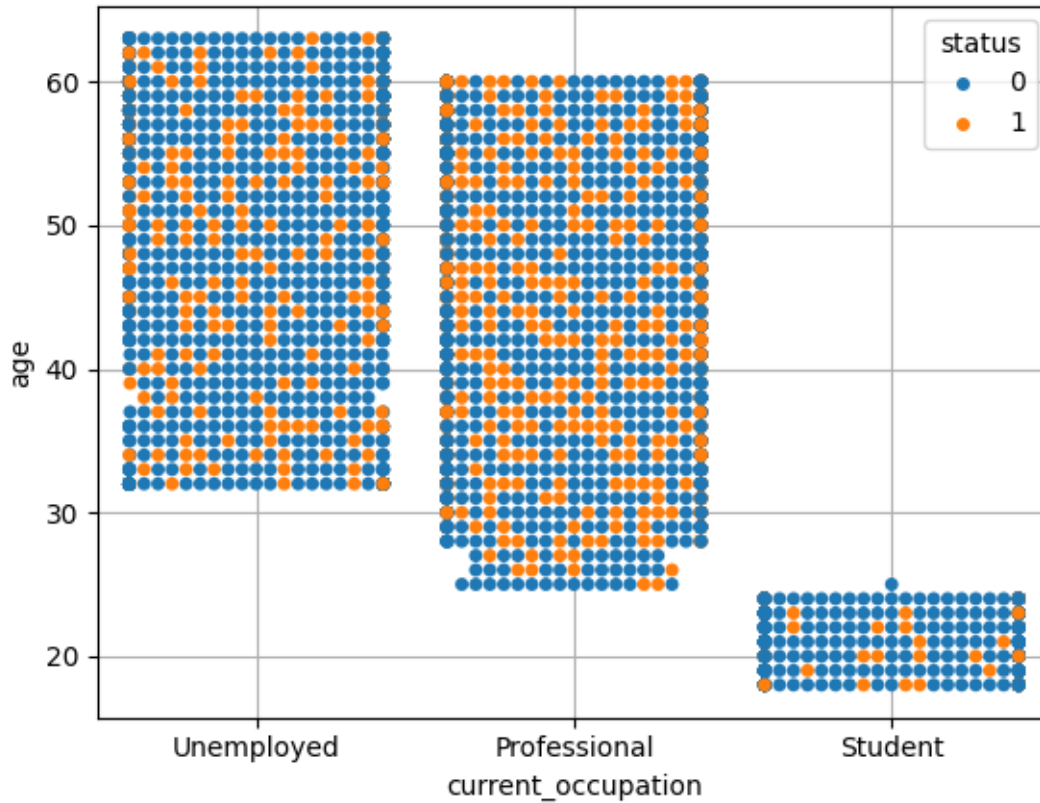


Observations on the relation between the Current Occupation and successful leads.

- It seems that students are unlikely to become paid customers (11.7%)
- About 27% of unemployed leads end up paying
- More than 35% of professionals (who happened to be the most common type of lead) end up becoming paid customers.

It will be therefore interesting evaluating if age can differentiate these 3 scenarios.

```
[51]: sns.swarmplot(data=data, x="current_occupation", y="age", hue="status")  
      plt.grid(visible=True, which="both", axis="both")  
      plt.show()
```

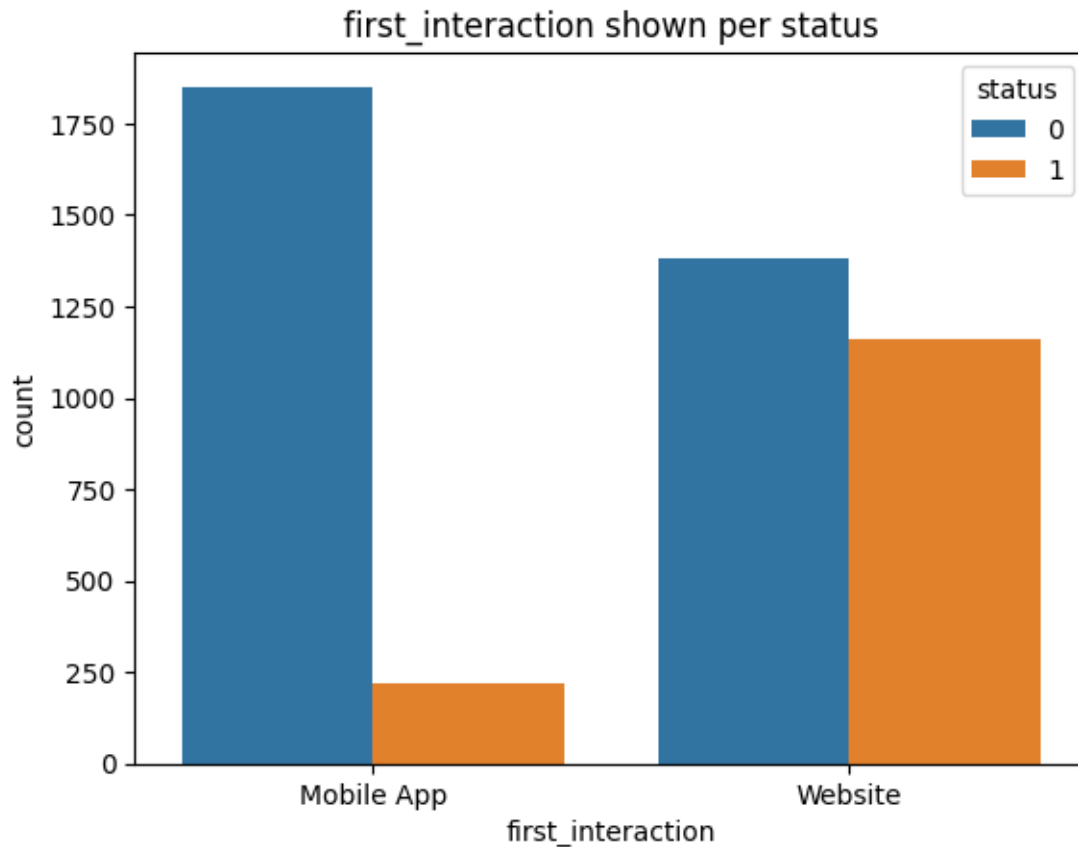


It seems that the age of a specific occupation doesn't play a big role. However, there is a clear distinction at 25 years old (by this age masters are usually finished. Probably PhDs are considered as Professionals.)

Let's take a look on the impact of the type of first impression.

```
[52]: barplot_hued(data, "first_interaction", hue="status")
```

		proportion	count
first_interaction	status		
Mobile App	0	89.46860	1852
	1	10.53140	218
Website	0	54.40598	1383
	1	45.59402	1159



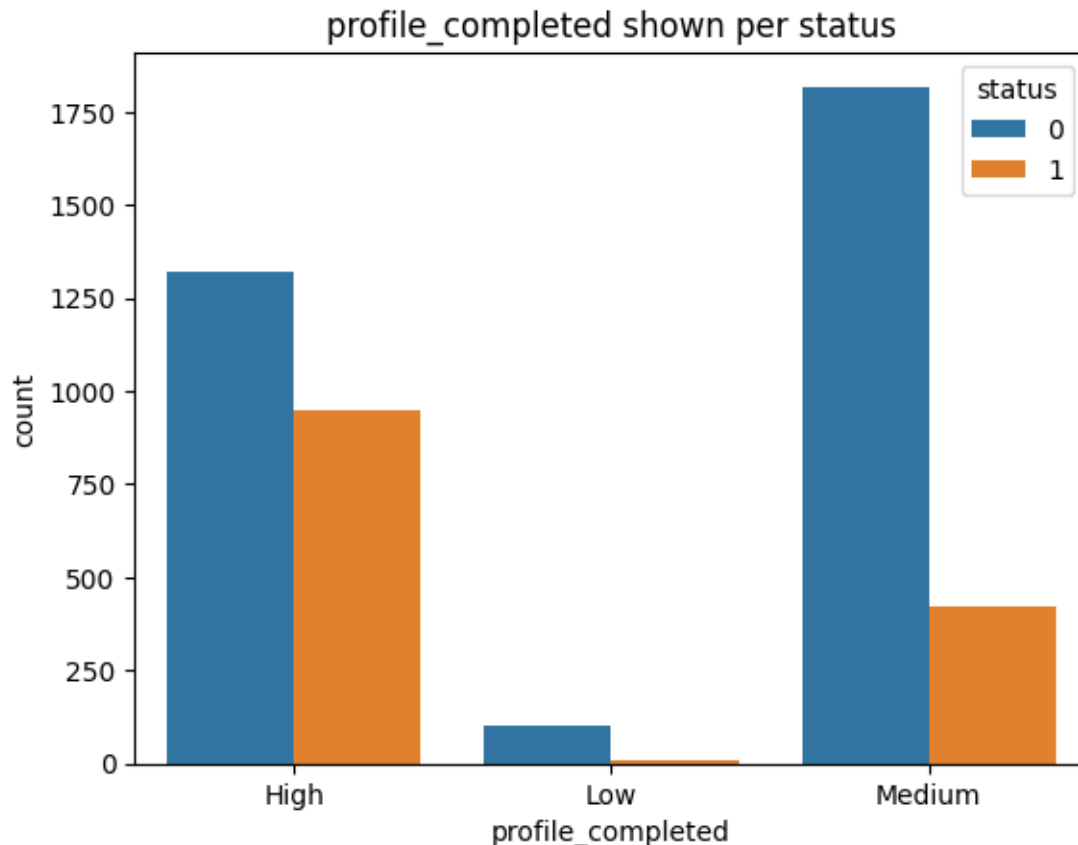
Observations on the relation between the type of First Interaction and successful leads.

- More people have the first impression through the Website
- **More than 45% of those who saw the Website first end up paying!**
- Only 10% of those who saw the Mobile App first end up paying.

Let's check if the level of Profile Completion plays a role.

```
[53]: barplot_hued(data, "profile_completed", hue="status")
```

		proportion	count
profile_completed	status		
High	0	58.21555	1318
	1	41.78445	946
Low	0	92.52336	99
	1	7.47664	8
Medium	0	81.12450	1818
	1	18.87550	423

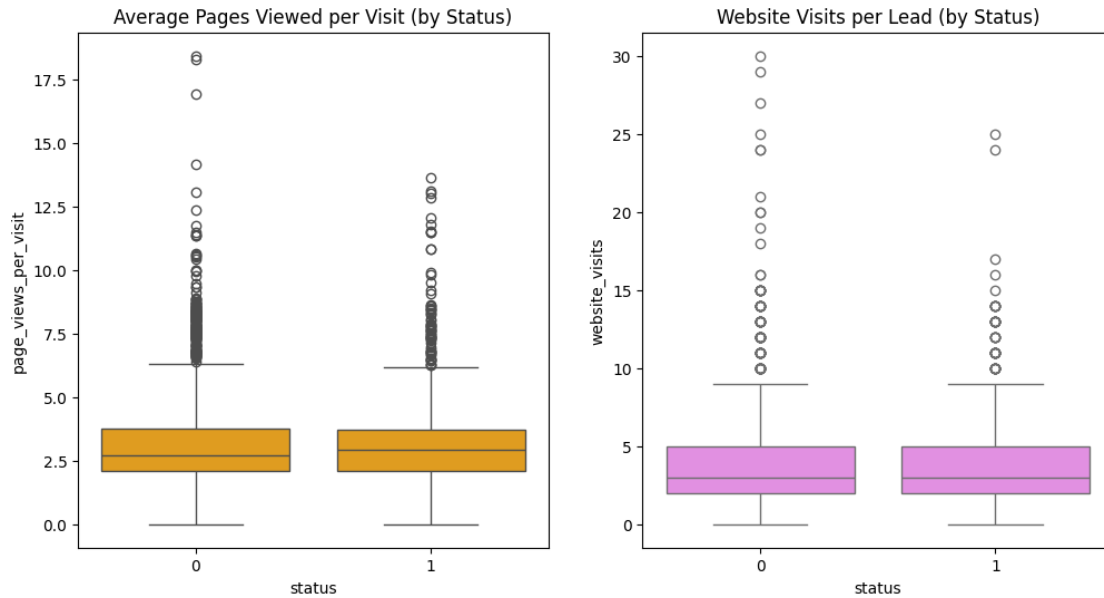


Observations on the relation between the level of Profile Completion and successful leads.

- Not only the absolute amount of paid customers increases with the level of profile completion, but also the relative amount of paid customers for a given level of profile completion increases (low_relative_proportion < medium_relative_proportion < high_relative_proportion <=> 7.5% < 18.9% < 41.8%).
- 4/10 leads who have a high level of profile completion become paid customers.

Let's check if the Pages Viewed and the amount of times a lead has Visited the Website are meaningful

```
[127]: fig, axs = plt.subplots(1, 2, figsize=(12, 6))
sns.boxplot(data=data, x="status", y="page_views_per_visit", color="orange", ax=axs[0])
sns.boxplot(data=data, x="status", y="website_visits", color="violet", ax=axs[1])
axs[0].set_title("Average Pages Viewed per Visit (by Status)")
axs[1].set_title("Website Visits per Lead (by Status)")
plt.show()
```



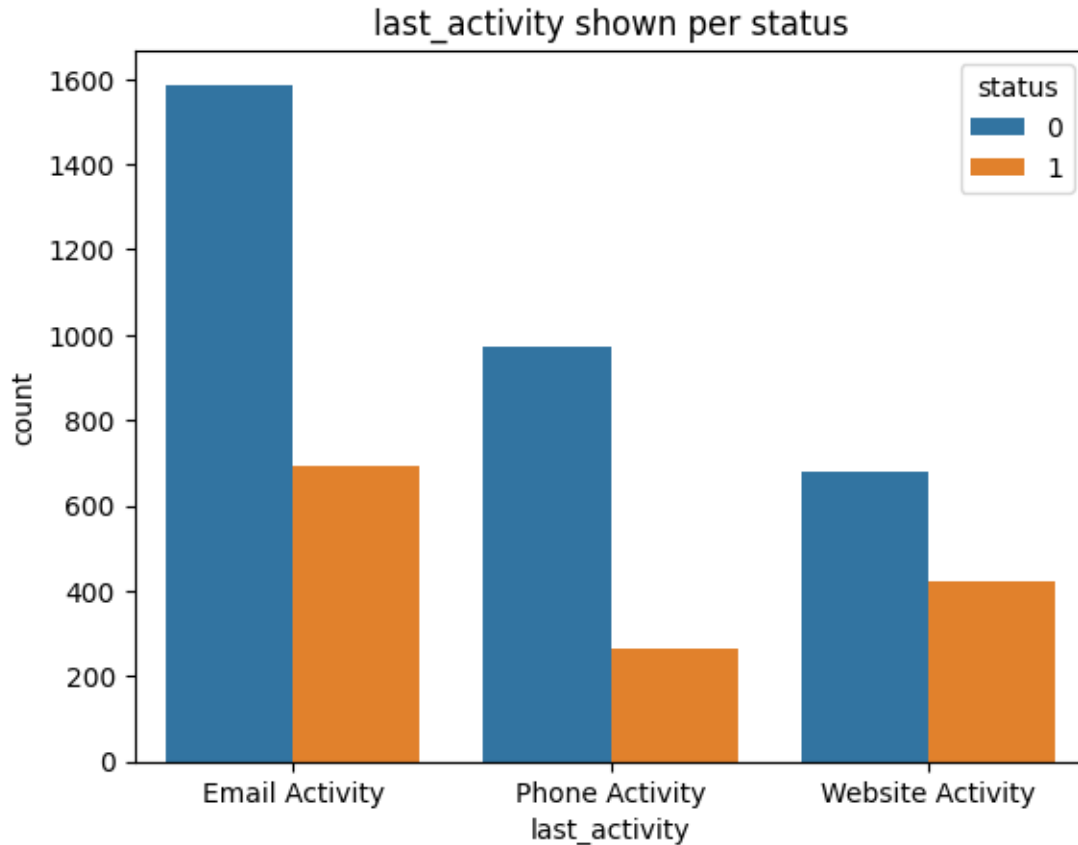
Observations regarding the Website visits, Average Page Viewed per visit and its relation to the successful customers:

- The boxplots for both categories are very similar respectively.
- The outliers are reasonable for both categories.
- From the heatmap it is known that the distributions of each category (hued by the Status) are very similar respectively.
- Overall, these are **apparently** not very useful categories.

Let's check if the Last Activity has an impact on the number of successful leads

```
[55]: barplot_hued(data, "last_activity", hue="status")
```

		proportion	count
last_activity	status		
Email Activity	0	69.66637	1587
	1	30.33363	691
Phone Activity	0	78.68720	971
	1	21.31280	263
Website Activity	0	61.54545	677
	1	38.45455	423



Observations on whether the type of Last Activity has an impact on paid customers

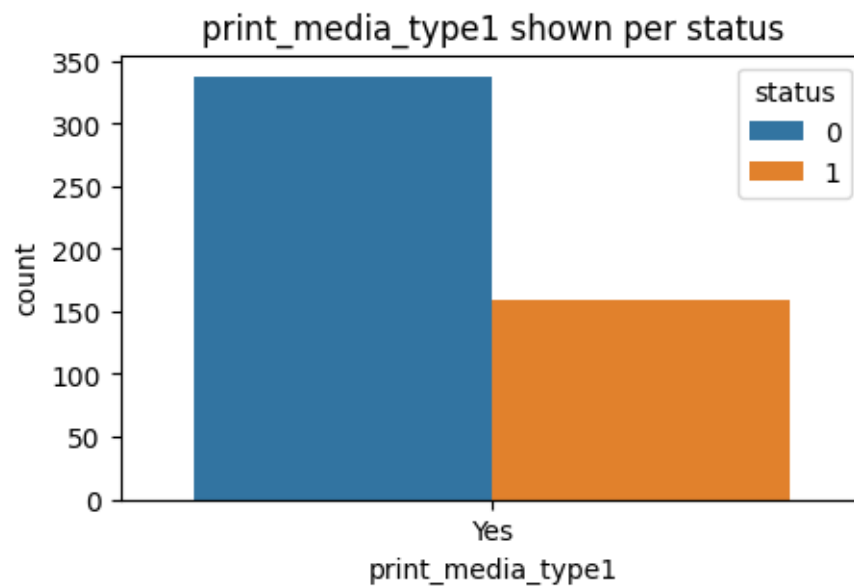
- Most customers had their last activity via Email. From those, 30% became paid customers.
- Many customers had their last activity via their Phone. From those, only 21% became paid customers.
- Many customers had their last activity via the Website. From those, 38% became paid customers.
- It is pertinent to mention that the absolute paid customers per type also matter. Namely, even if it is known that if the last activity was through the Website, then statistically the leads are more likely to become customers, it is important to consider that from Email 691 leads became customers, from Website 423 became customers, and from a Phone activity only 263.
- If a customer had the last activity through a phone activity, then both in relative and absolute numbers, it is less likely that the lead will become a paid customer.

Let's find what is the best channel in terms of the (highest) lead conversion rate

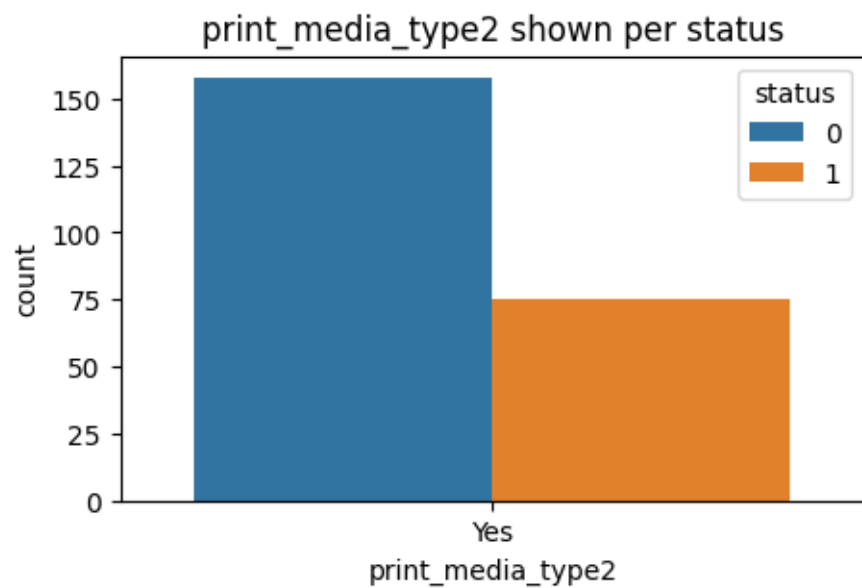
Let's start with the channel of Newspaper

```
[56]: channel_types = ["print_media_type1", "print_media_type2", "digital_media",
    ↪ "educational_channels", "referral"]
    for channel in channel_types:
        barplot_hued(data[data[channel] == "Yes"], channel, hue="status",
    ↪ figsize=(5,3))
        print("-"*60 + "\n")
```

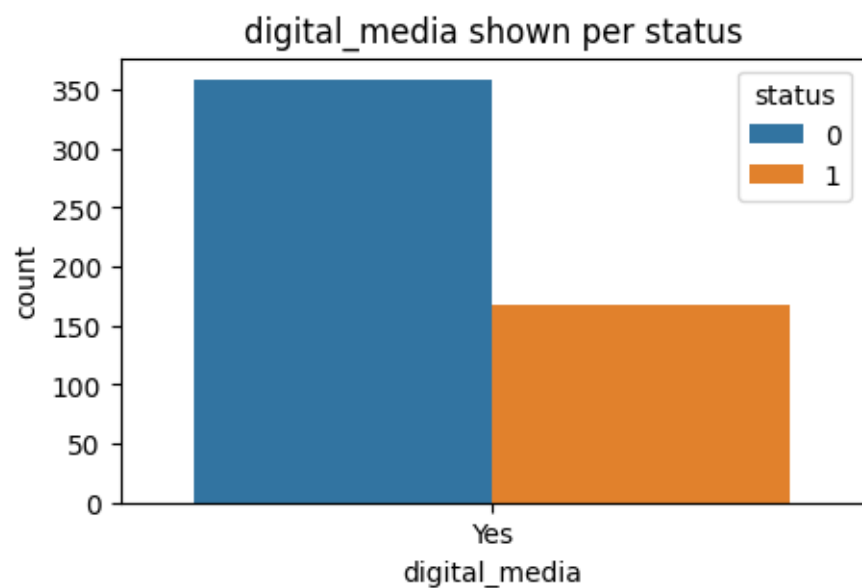
		proportion	count
print_media_type1	status		
Yes	0	68.00805	338
	1	31.99195	159



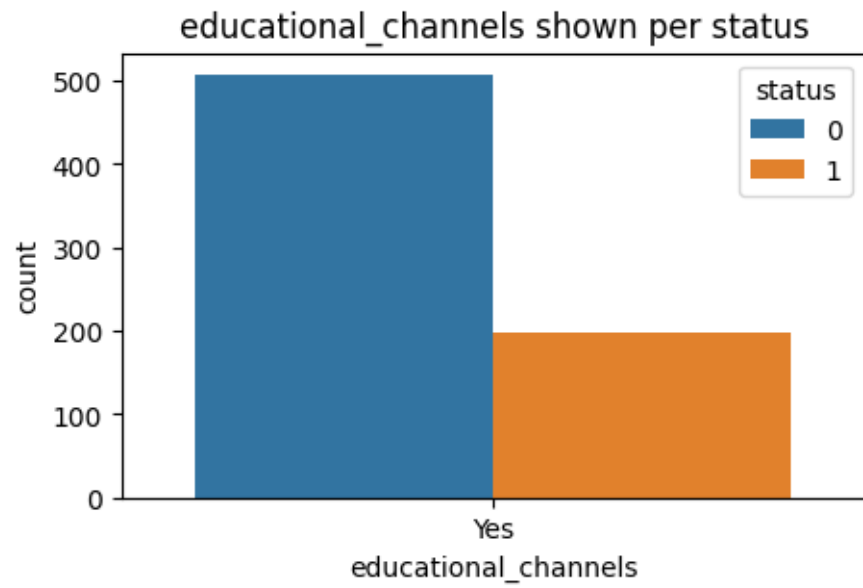
		proportion	count
print_media_type2	status		
Yes	0	67.81116	158
	1	32.18884	75



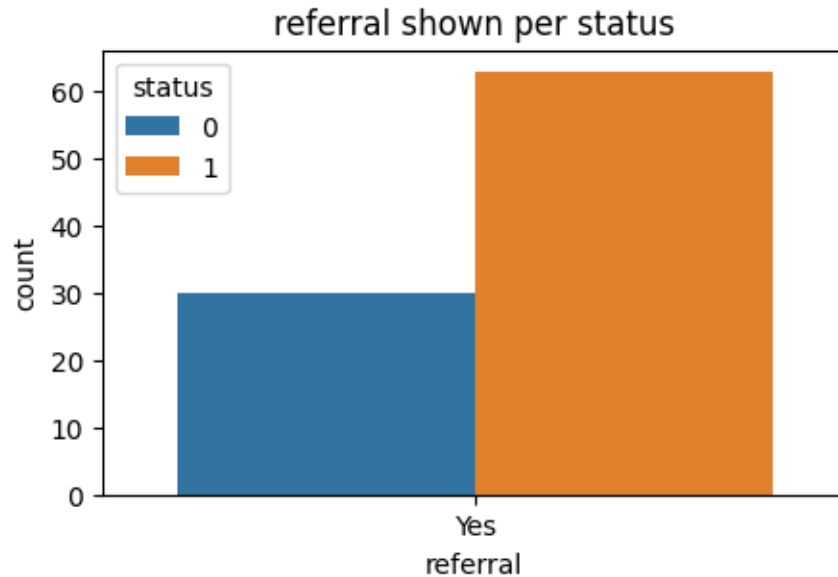
		proportion	count
digital_media status			
Yes	0	68.12144	359
	1	31.87856	168



		proportion	count
educational_channels	status		
Yes	0	72.05674	508
	1	27.94326	197



		proportion	count
referral	status		
Yes	0	32.25806	30
	1	67.74194	63



Observations on the channel used to get a lead

- Newspapers brought 159 paid customers representing 31.9% of those leads acquired by this channel.
- The Magazine brought 75 paid customers representing 32.2% of those leads acquired by this channel.
- Digital Platforms brought 168 paid customers representing 31.9% of those leads acquired by this channel.
- Educational Channels brought 197 paid customers representing 27.9% of those leads acquired by this channel.
- Newspaper brought 63 paid customers representing 67.7% of those leads acquired by this channel.
- **The highest lead conversion rate is clearly from Referrals.** In fact, referral's lead conversion rate is **more than twice** that of the second best channel.
- Educational Channels, Digital Platforms and Newspapers bring similar amount of paid customers.
- Magazines bring too few paid customers in comparison to the rest.

1.8 Building a Decision Tree model

A Decision Tree offers a model that works like a flowchart. Every node represents a question, and each subnode represents an “answer” to the question. The goal is to find the right questions to effectively and efficiently divide the dataset.

```
[57]: X = data.drop(["status"], axis=1)
      Y = data["status"]

      X = pd.get_dummies(X, drop_first=True)

      # Splitting the dataset into train (80%) and test (20%) datasets.
      x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.2,
      ↪shuffle = True, random_state = 1)
```

```
[58]: print("Shape of Training set : ", x_train.shape)
      print("Shape of test set : ", x_test.shape)
      print("Percentage of classes in training set:") # Classes being 0 and 1
      print(y_train.value_counts(normalize=True)) # How are the values relatively
      ↪distributed
      print("Percentage of classes in test set:")
      print(y_test.value_counts(normalize=True))
```

```
Shape of Training set : (3689, 16)
Shape of test set : (923, 16)
Percentage of classes in training set:
status
0    0.70480
1    0.29520
Name: proportion, dtype: float64
Percentage of classes in test set:
status
0    0.68797
1    0.31203
Name: proportion, dtype: float64
```

1.8.1 Understanding what would be a good model

We want a model that can predict which leads are more likely to become paid customers so that the economic and human efforts are optimized, and at the same time, generate more money.

All Scenarios

- **True Positive:** a lead is **predicted** as someone that **would become a paid customer**, actions are taken, and **this lead indeed becomes a paid customer**.
- **True Negative:** a lead is **predicted** as someone that **would not become a paid customer**, no actions are taken, and **this lead indeed does not become a paid customer**.
- **False Positive:** a lead is **predicted** as someone that **would become a paid customer**, actions are taken, and in fact this lead was incorrectly labeled because, **this lead will not become a paid customer**.
- **False Negative:** a lead is **predicted** as someone that **would not become a paid customer**, no actions are taken, and in fact this lead was incorrectly labeled because, **this lead would have become a paid customer**.

What is worse? Losing a potential customer because the prediction was wrong! The False Negatives should be reduced. Recall should be maximized! $\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$

A high Recall means that most of the positives (leads that will become paid customers) are indeed predicted as such. A high Recall implies a low number of False Negatives.

Let's define a function that gives a lot of metrics to evaluate the performance of a model

```
[84]: # Function to print the classification report and get confusion matrix in a
      ↪ proper format

def metrics_score(actual, predicted):
    # Prints the classification report, built-in with sklearn
    print(classification_report(actual, predicted))
    # Generates a confusion matrix
    cm = confusion_matrix(actual, predicted)

    # Figure size
    plt.figure(figsize = (8, 5))
    # Heatmap
    sns.heatmap(cm, annot = True, fmt = '.2f', xticklabels = ['Not Converted',
    ↪ 'Converted'], yticklabels = ['Not Converted', 'Converted'])
    plt.ylabel('Actual')
    plt.xlabel('Predicted')
    plt.show()
```

1.8.2 Building the Decision Tree Model

The sklearn.tree function *DecisionTreeClassifier* will be used (because the output is a discrete value).

```
[60]: d_tree = DecisionTreeClassifier(random_state=1)

      # Fitting the model with the training set
      d_tree.fit(x_train, y_train)
```

```
[60]: DecisionTreeClassifier(random_state=1)
```

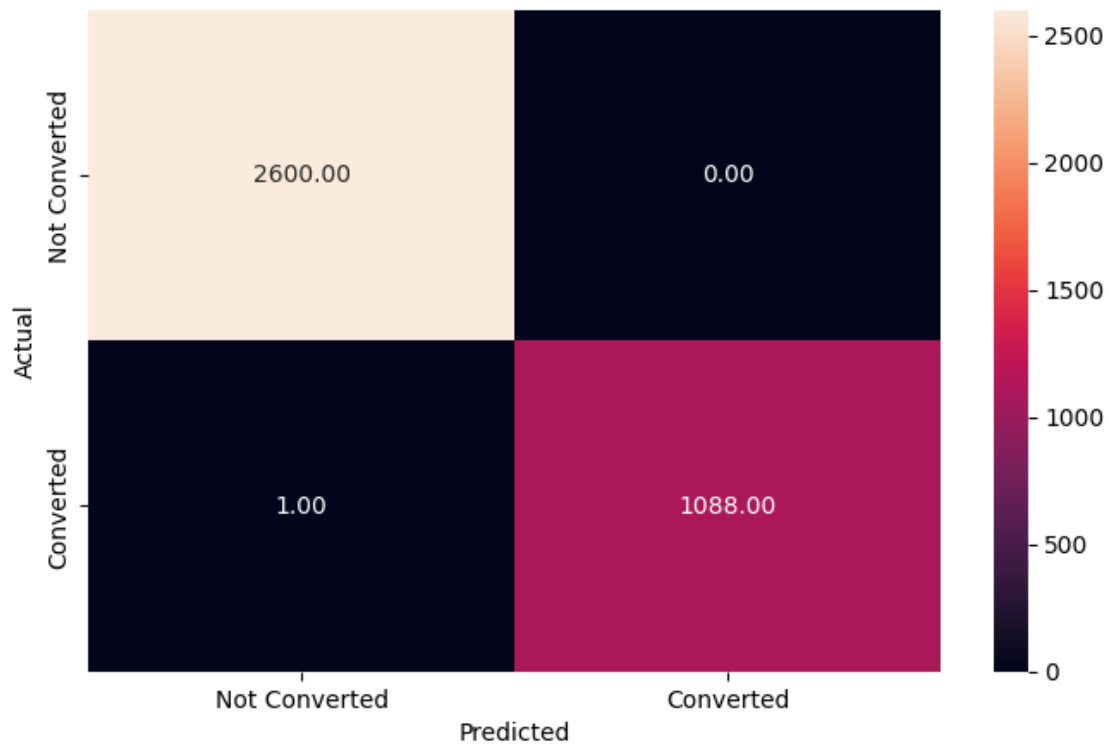
Let's check the performance of the model on the training data

```
[61]: # Checking performance on the training data
      y_pred_train1 = d_tree.predict(x_train)

      metrics_score(y_train, y_pred_train1)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2600

	1	1.00	1.00	1.00	1089
accuracy				1.00	3689
macro avg		1.00	1.00	1.00	3689
weighted avg		1.00	1.00	1.00	3689



The model fits really well the training data. There is only one False Negative.

Let's see how the model performs on the test data

```
[62]: # Checking performance on the test data
y_pred_test1 = d_tree.predict(x_test)

metrics_score(y_test, y_pred_test1)
```

	precision	recall	f1-score	support
0	0.86	0.87	0.87	635
1	0.71	0.70	0.71	288
accuracy			0.82	923
macro avg	0.79	0.79	0.79	923

weighted avg 0.82 0.82 0.82 923



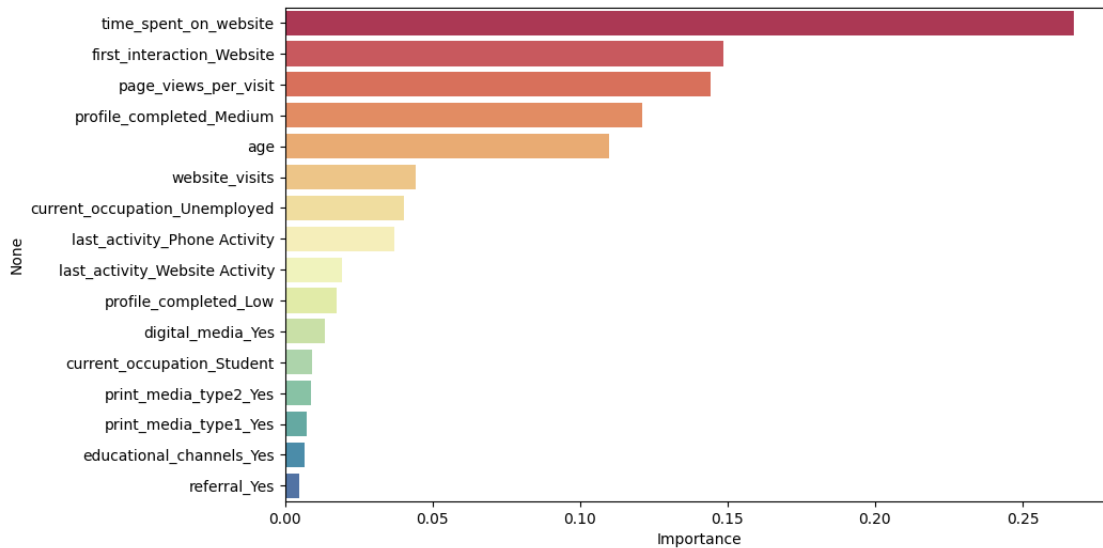
Observations on the model when fitted the test data:

- All the metrics show that the model is performing worse with the test dataset. Moreover, there is a Recall of 70%, which means that 3 out of every 10 leads that could become paid customers are incorrectly labeled.
- The model is overfitting the training dataset.

Let's calculate a measure of relevance per feature called importance.

```
[85]: # Obtain the importance of each feature in the decision tree
importances = d_tree.feature_importances_
# Get the name of all the columns that could help the model
columns = X.columns
# Generate a Dataframe for plotting purposes in descending order
importance_df = pd.DataFrame(importances, index = columns, columns = ['Importance']).sort_values(by = 'Importance', ascending = False)
# Barplot: Importances of features
plt.figure(figsize = (10, 6))
sns.barplot(data = importance_df, x = importance_df.Importance, y = importance_df.index, palette="Spectral")
```

[85]: <Axes: xlabel='Importance', ylabel='None'>

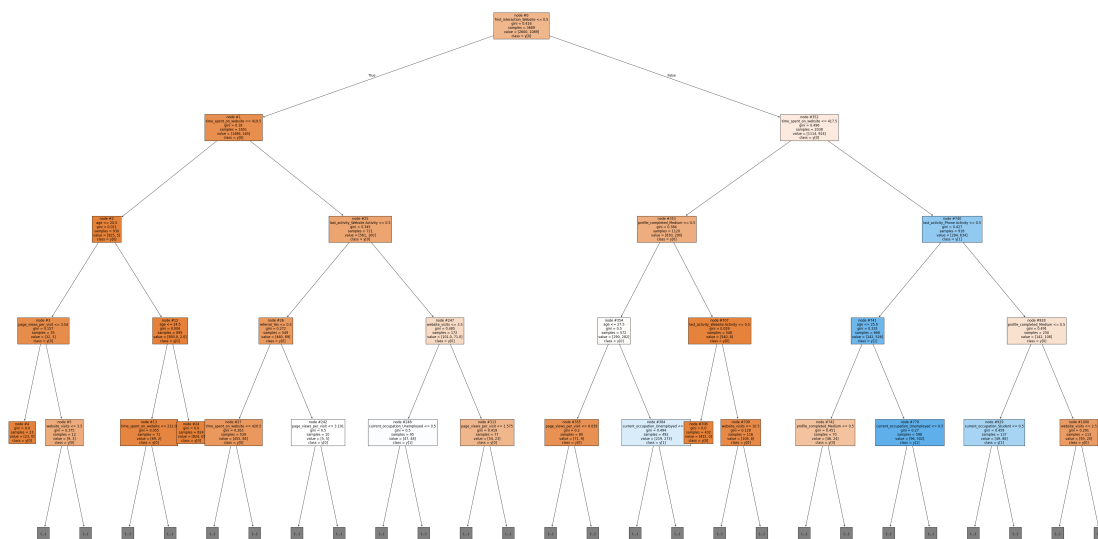


Observations:

- As previously inferred, the time spent on the website is one of the features with the most importance.
- Even if Referrals was the best channel, it does not have much importance. In fact, other channels such as digital media are considered more important. This is why it is recommended to avoid jumping to conclusions too early about the relevance of features.
- Having the first interaction through the website, the amount of pages viewed per visit, having a medium level of profile completion and the age of the lead are the features with most importance in this rudimentary decision tree model.
- There seems to be a lot of features considered with little to no importance. Therefore, it might be necessary to prune the tree. ##### Let's visualize the tree

```
[97]: # The features of the tree are formatted into a list.
features = list(X.columns)

# The tree is plotted. A depth of 4 was chosen to enhance readability.
plt.figure(figsize = (73, 40))
tree.plot_tree(d_tree, max_depth = 4, feature_names = features, filled = True,
    ↳fontsize = 12, node_ids = True, class_names = True)
plt.show()
```



Observations:

- The tree needs to be pruned. There are a lot of branches with very little variation. Moreover, the depth is more than 5 and it is not very readable.

1.8.3 Let's see if the tree can be shortened by tuning of parameters.

Tuning hyperparameters with GridSearchCV for Decision Tree

```
[65]: # Choose the type of classifier
d_tree_tuned = DecisionTreeClassifier(random_state = 7, class_weight = {0: 0.3,
↪1: 0.7})

# Grid of parameters to choose from
parameters = {'max_depth': np.arange(2, 10),
              'criterion': ['gini', 'entropy'],
              'min_samples_leaf': [5, 10, 20, 25]
              }

# Type of scoring used to compare parameter combinations - recall score for
↪class 1
scorer = metrics.make_scorer(recall_score, pos_label = 1)

# Run the grid search
grid_obj = GridSearchCV(d_tree_tuned, parameters, scoring = scorer, cv = 5)
grid_obj = grid_obj.fit(x_train, y_train)

# Set the classifier to the best combination of parameters
```

```
d_tree_tuned = grid_obj.best_estimator_

# Fit the best algorithm to the data
d_tree_tuned.fit(x_train, y_train)
```

```
[65]: DecisionTreeClassifier(class_weight={0: 0.3, 1: 0.7}, criterion='entropy',
                             max_depth=3, min_samples_leaf=5, random_state=7)
```

Let's check the performance of the model on the training data

```
[66]: # Checking performance on the training data
y_pred_train1_tuned = d_tree_tuned.predict(x_train)

metrics_score(y_train, y_pred_train1_tuned)
```

	precision	recall	f1-score	support
0	0.95	0.74	0.83	2600
1	0.59	0.91	0.72	1089
accuracy			0.79	3689
macro avg	0.77	0.82	0.77	3689
weighted avg	0.84	0.79	0.80	3689



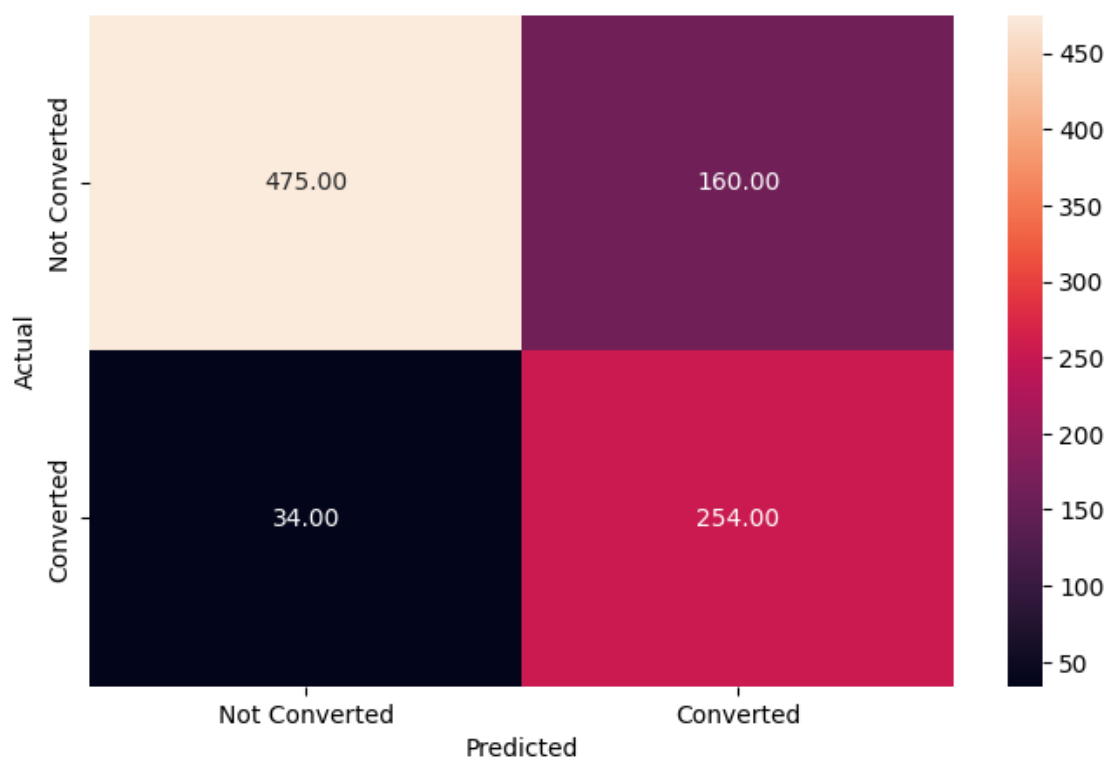
The model has a recall of 91% on the training data.

Let's see how it performs on the test set

```
[98]: # Checking performance on the test data
y_pred_test1_tuned = d_tree_tuned.predict(x_test)

metrics_score(y_test, y_pred_test1_tuned)
```

	precision	recall	f1-score	support
0	0.93	0.75	0.83	635
1	0.61	0.88	0.72	288
accuracy			0.79	923
macro avg	0.77	0.81	0.78	923
weighted avg	0.83	0.79	0.80	923

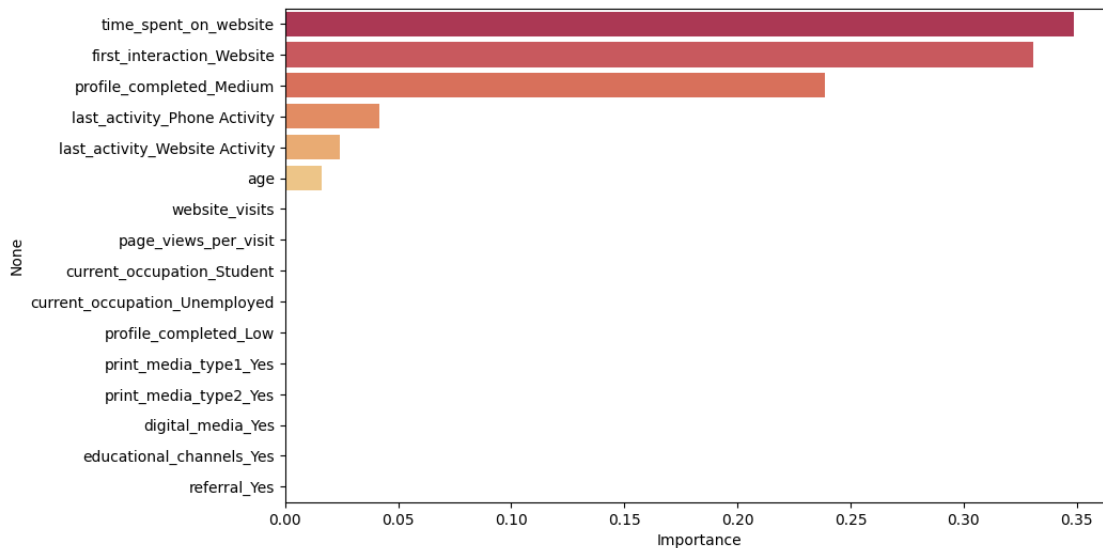


Observations:

- The Recall is 88%. This is only 3% less than the one from the training dataset. There is not much overfitting.
- The recall of the tuned decision tree is way higher than that of the regular decision tree (88% > 70%). GridSearchCV was indeed effective. ##### Let's calculate a measure of relevance per feature called importance.

```
[101]: # Obtain the importance of each feature in the decision tree
importances = d_tree_tuned.feature_importances_
# Get the name of all the columns that could help the model
columns = X.columns
# Generate a Dataframe for plotting purposes in descending order
importance_df = pd.DataFrame(importances, index = columns, columns = ['Importance']).sort_values(by = 'Importance', ascending = False)
# Barplot: Importances of features
plt.figure(figsize = (10, 6))
sns.barplot(data = importance_df, x = importance_df.Importance, y = importance_df.index, palette="Spectral")
```

```
[101]: <Axes: xlabel='Importance', ylabel='None'>
```



Observations:

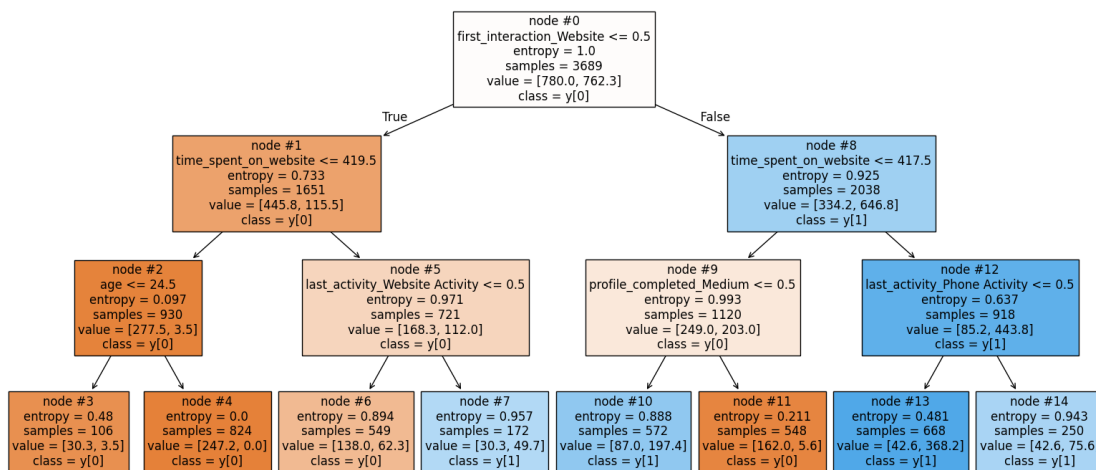
- As previously inferred, the time spent on the website is still one of the features with the most importance.
- Even if the amount of pages viewed per visit was very relevant for the rudimentary decision tree, it has no importance in the tuned decision tree. On the other hand, having the last activity through a phonecall/SMS or the website were previously irrelevant and now they are relevant.
- Having the first interaction through the website, having a medium level of profile completion,

last interaction through a phonecall/SMS or website and the age of the lead are the features with most importance in this tuned decision tree model.

- The decision tree has been pruned. ##### Let's visualize the Tuned Decision Tree

```
[105]: # The features of the tree are formatted into a list.
features = list(X.columns)

# The tree is plotted. No max depth is selected.
plt.figure(figsize = (20, 9))
tree.plot_tree(d_tree_tuned, feature_names = features, filled = True, fontsize=
↳ 12, node_ids = True, class_names = True)
plt.show()
```



Observations:

- The Tuned Decision Tree has 88% Recall, reducing the number of false negatives while keeping a rather straightforward flowchart with a depth of 4.
- **The leftmost and rightmost branches can be pruned** because they are not bringing new information.

1.9 Random Forest

A (Random) Forest is made by multiple (Decision) Trees, hence the name. It is an ensemble method, with the drawback of losing interpretability and readability. ##### Let's build a Random Forest Model to see if it performs better than the Tuned Decision Tree.

```
[107]: # Fitting the Random Forest classifier on the training data
rf_estimator = RandomForestClassifier(class_weight = "balanced", random_state =
↳ 1)

rf_estimator.fit(x_train, y_train)
```

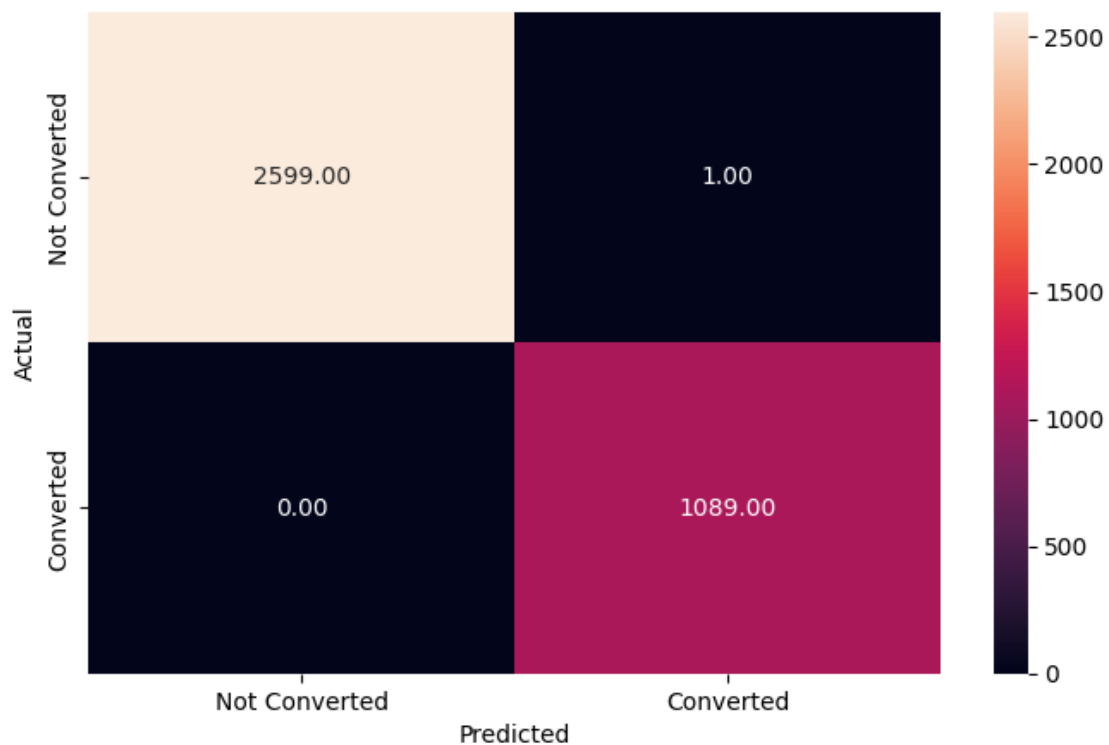
```
[107]: RandomForestClassifier(class_weight='balanced', random_state=1)
```

Let's check the performance of the model on the training data

```
[108]: # Checking performance on the training data
y_pred_train1_rf = rf_estimator.predict(x_train)

metrics_score(y_train, y_pred_train1_rf)
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2600
1	1.00	1.00	1.00	1089
accuracy			1.00	3689
macro avg	1.00	1.00	1.00	3689
weighted avg	1.00	1.00	1.00	3689



The model fits really well the training data. There are no False Negative.

Let's see how the model performs on the test data


```
[109]: # Checking performance on the test data
y_pred_test1_rf = rf_estimator.predict(x_test)

metrics_score(y_test, y_pred_test1_rf)
```

	precision	recall	f1-score	support
0	0.88	0.94	0.91	635
1	0.84	0.71	0.77	288
accuracy			0.87	923
macro avg	0.86	0.83	0.84	923
weighted avg	0.87	0.87	0.86	923



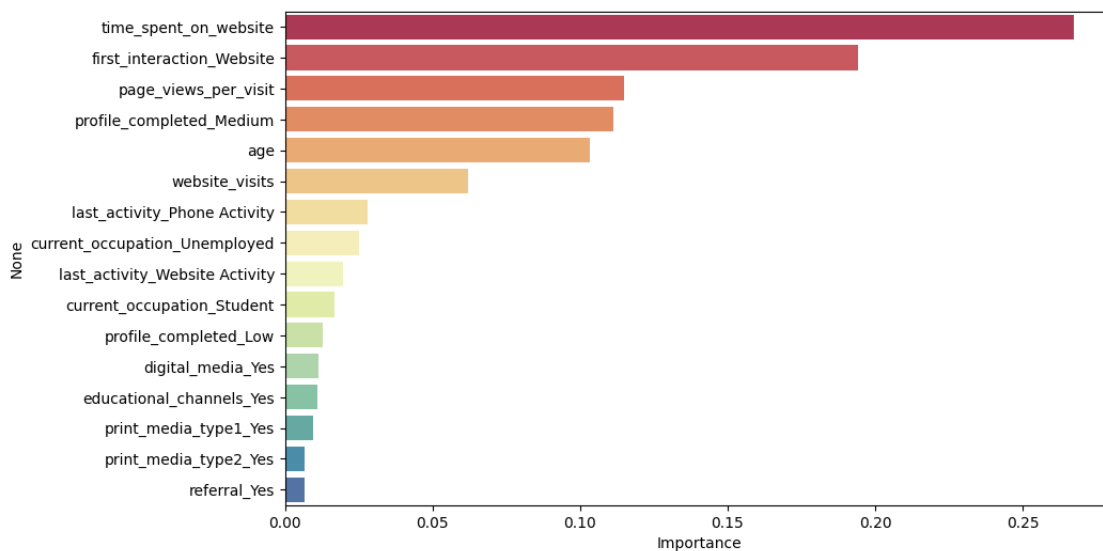
Observations on the model when fitted the test data:

- All the metrics show that the model is performing worse with the test dataset. Moreover, there is a Recall of 71%, which means that approximately 3 out of every 10 leads that could become paid customers are incorrectly labeled.
- The model is overfitting the training dataset.

Let's calculate a measure of relevance per feature called importance.

```
[110]: # Obtain the importance of each feature in the random forest
importances = rf_estimator.feature_importances_
# Get the name of all the columns that could help the model
columns = X.columns
# Generate a Dataframe for plotting purposes in descending order
importance_df = pd.DataFrame(importances, index = columns, columns = ['Importance']).sort_values(by = 'Importance', ascending = False)
# Barplot: Importances of features
plt.figure(figsize = (10, 6))
sns.barplot(data = importance_df, x = importance_df.Importance, y = importance_df.index, palette="Spectral")
```

```
[110]: <Axes: xlabel='Importance', ylabel='None'>
```



Observations:

- The time spent on the website is the feature with the most importance.
- Having the first interaction through the website, the amount of pages viewed per visit, having a medium level of profile completion, the age of the lead and the amount of website visits are the features with most importance in this rudimentary random forest model.
- There seems to be a lot of features considered with little to no importance. Therefore, it might be necessary to prune the tree. ##### Let's consider the first tree in this (random) forest

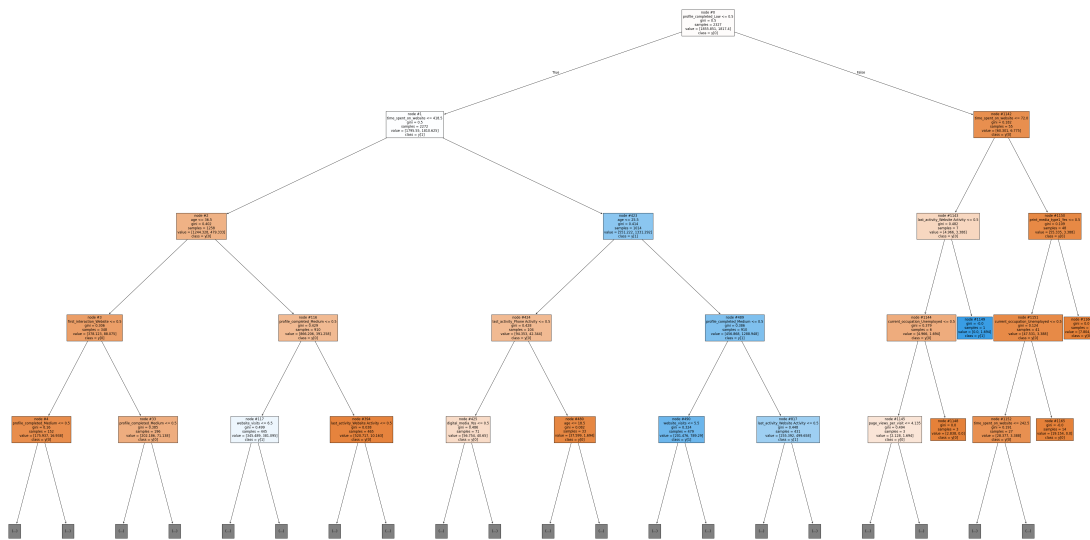
```
[115]: print(f"There are {rf_estimator.n_estimators} trees in this (random) forest")
first_tree_rf = rf_estimator[0]
```

There are 100 trees in this (random) forest

Let's visualize the first of all the trees in this forest.

```
[116]: # The features of the tree are formatted into a list.
features = list(X.columns)

# The tree is plotted. A depth of 4 was chosen to enhance readability.
plt.figure(figsize = (73, 40))
tree.plot_tree(first_tree_rf, max_depth = 4, feature_names = features, filled =
↳ True, fontsize = 12, node_ids = True, class_names = True)
plt.show()
```



Observations:

- The tree needs to be pruned. There a lot of branches with very little variation. Moreover, the depth is more than 5, which reduces readability.

1.9.1 Let's see if the Random Forest model can be improved by tuning the parameters.

Tuning hyperparameters with GridSearchCV for Random Forest

```
[118]: # Choose the type of classifier
rf_estimator_tuned = RandomForestClassifier(criterion = "entropy", random_state
↳ 1)

# Grid of parameters to choose from
parameters_rf = {"n_estimators": [110, 120],
"max_depth": [6, 7],
"min_samples_leaf": [20, 25],
```

```

    "max_features": [0.8, 0.9],
    "max_samples": [0.9, 1],
    "class_weight": ["balanced",{0: 0.3, 1: 0.7}]
    }

# Type of scoring used to compare parameter combinations - recall score for
↪ class 1
scorer = metrics.make_scorer(recall_score, pos_label = 1)

# Run the grid search
grid_obj = GridSearchCV(rf_estimator_tuned, parameters_rf, scoring = scorer, cv
↪= 5)

grid_obj = grid_obj.fit(x_train, y_train)

# Set the classifier to the best combination of parameters
rf_estimator_tuned = grid_obj.best_estimator_

```

Fitting the data into this tuned random forest

```

[119]: rf_estimator_tuned.fit(x_train, y_train)

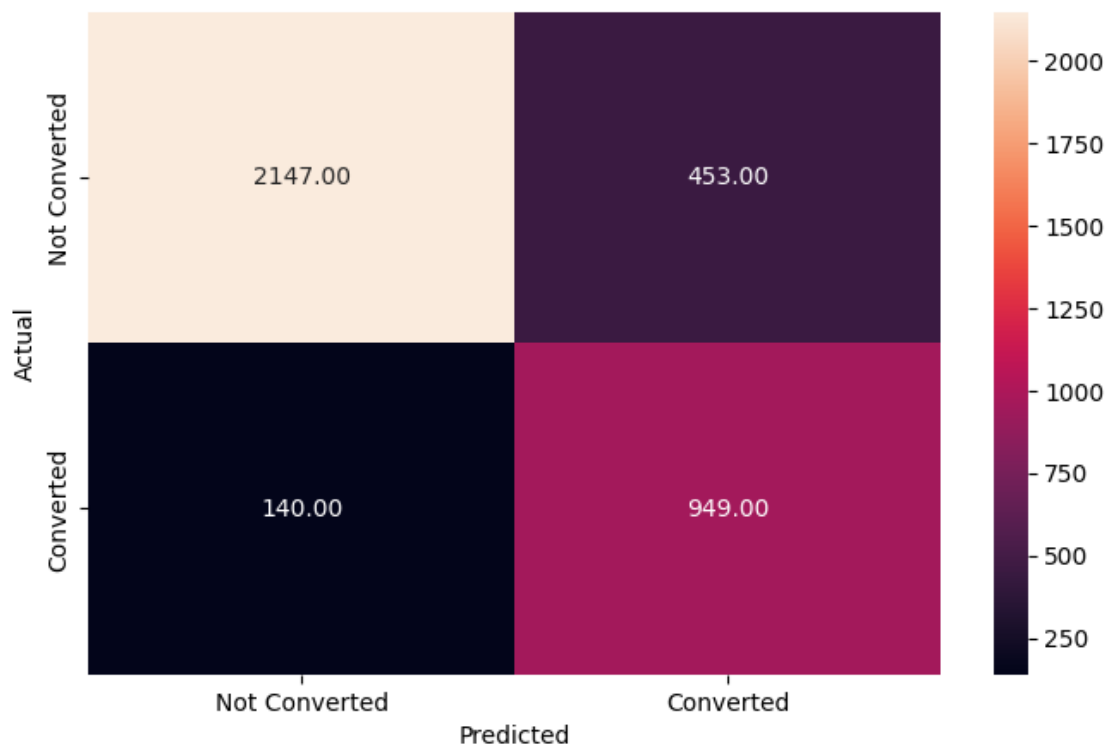
[119]: RandomForestClassifier(class_weight='balanced', criterion='entropy',
                             max_depth=6, max_features=0.8, max_samples=0.9,
                             min_samples_leaf=25, n_estimators=120, random_state=1)

[120]: # Checking performance on the training data
y_pred_train1_rf_tuned = rf_estimator_tuned.predict(x_train)

metrics_score(y_train, y_pred_train1_rf_tuned)

```

	precision	recall	f1-score	support
0	0.94	0.83	0.88	2600
1	0.68	0.87	0.76	1089
accuracy			0.84	3689
macro avg	0.81	0.85	0.82	3689
weighted avg	0.86	0.84	0.84	3689



The model has a recall of 87% on the training data.

Let's see how it performs on the test set

```
[123]: # Checking performance on the test data
y_pred_test1_rf_tuned = rf_estimator_tuned.predict(x_test)

metrics_score(y_test, y_pred_test1_rf_tuned)
```

	precision	recall	f1-score	support
0	0.92	0.84	0.88	635
1	0.70	0.85	0.77	288
accuracy			0.84	923
macro avg	0.81	0.84	0.82	923
weighted avg	0.85	0.84	0.84	923

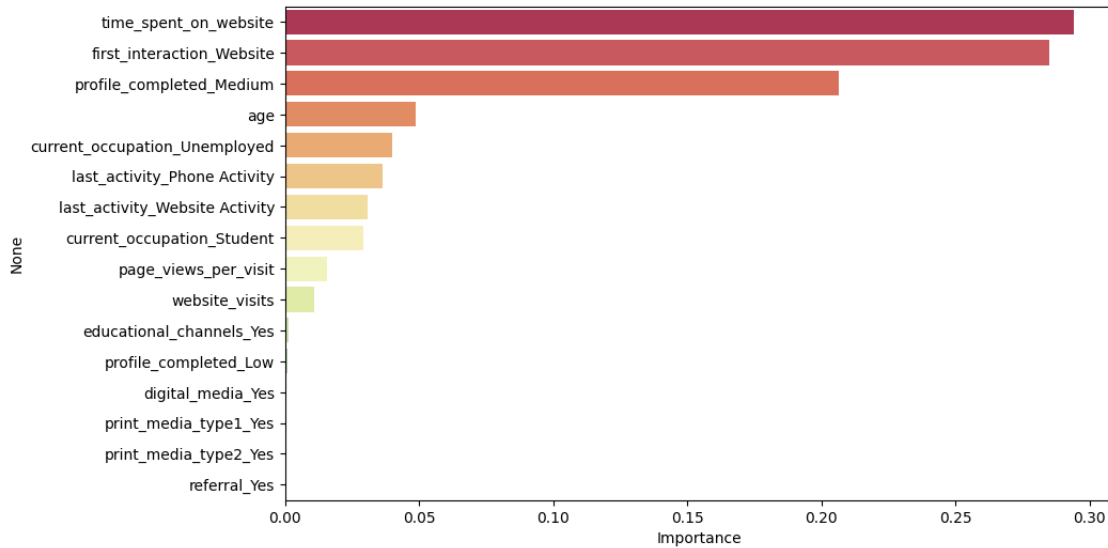


Observations:

- The Recall is 85%. This is only 2% less than the one from the training dataset. There is almost no overfitting.
- The recall of the tuned decision tree is way higher than that of the regular decision tree (85% > 71%). GridSearchCV was indeed effective. ##### Let's calculate a measure of relevance per feature called importance.

```
[124]: # Obtain the importance of each feature in the decision tree
importances = rf_estimator_tuned.feature_importances_
# Get the name of all the columns that could help the model
columns = X.columns
# Generate a Dataframe for plotting purposes in descending order
importance_df = pd.DataFrame(importances, index = columns, columns = ['Importance'])
importance_df.sort_values(by = 'Importance', ascending = False)
# Barplot: Importances of features
plt.figure(figsize = (10, 6))
sns.barplot(data = importance_df, x = importance_df.Importance, y = importance_df.index, palette="Spectral")
```

```
[124]: <Axes: xlabel='Importance', ylabel='None'>
```



Observations:

- The time spent on the website is the feature with the most importance. Having the first interaction through the website and having a medium level of profile completion are all very important for this tuned random forest model.
- In this case, variables such as occupation are considered, meaning that the tuned random forest model takes into account more factors.
- Age, the type of the last activity and the pages viewed per visit are also important.

1.10 Actionable Insights and Recommendations

- More than 50% of the leads who become buyers are professionals. Indeed, in a professional world where soft skills are valued, it is logical for professionals to seek online education. It would be smart to target these people (older than 25).
- The first impression via the website yields way more paid customers than via the mobile app. In fact, 45% of those who first interact with the website end up paying. Therefore, it is recommended for the marketing department to promote ads that link to the website, and not the app.
- There is a 40% chance that a lead with a high profile completion ends up paying. Hence, emails suggesting to continue their profile completion might be helpful. What's more, 1/3 of leads whose last activity was through an email become leads.
- If the last activity was via a phone, only 1/5 turns into a paid customer. Therefore, the funds for phone communication should be low.
- Referrals have the highest conversion rate by a considerable margin. Therefore, giving incentives to current customers for promoting the program is a great idea.
- The best model (and yet not the most computationally expensive) is the Tuned Decision Tree. Only 12% of the potential customers will be inconveniently ignored. Moreover, the tree has depth 4, and can be pruned, enhancing readability.

- The model considers that the time spent on website, having a first interaction via the website, and having a medium profile completion are the most important variables.
- Assuring a great website experience is fundamental for the lead conversion. Completing a profile should be as user-friendly as possible.